

Towards deep learning with segregated dendrites

Guergiev^{1,2}, Lillicrap³ & Richards^{1,2,4}, eLife 2017

¹Department of Biological Sciences, University of Toronto Scarborough, Toronto, Canada;

²Department of Cell and Systems Biology, University of Toronto, Toronto, Canada;

³DeepMind, London, United Kingdom;

⁴Learning in Machines and Brains Program, Canadian Institute for Advanced Research, Toronto, Canada

Jesse Geerts

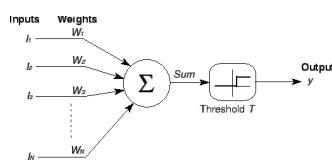
Computational Hippocampus Journal Club

31 Jan 2018

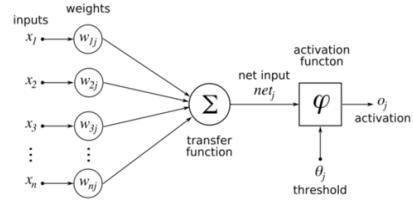
Outline

- **Introduction: comparing the brain to a neural network**
- The credit assignment problem in learning with networks
- Primer on the machine learning solution: back-propagation
- Biology's solution? Deep learning with segregated dendrites
- Results
- Discussion

Background: are artificial neural networks like the brain?

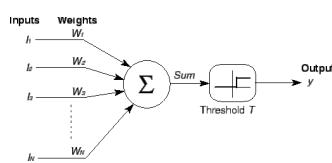


McCulloch & Pitts (1943):
"A logical calculus of the
ideas immanent in nervous
activity"

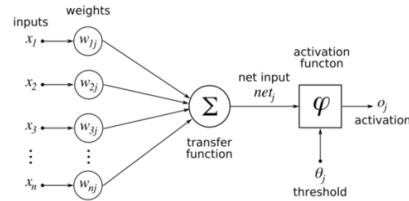


Rosenblatt (1958)
The perceptron

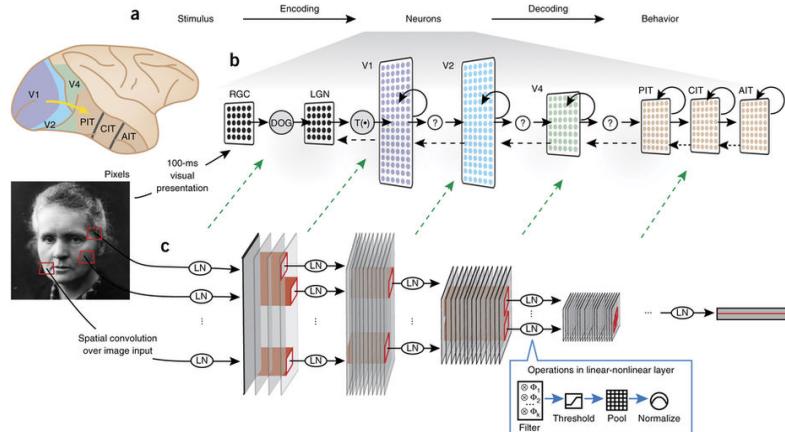
Background: are artificial neural networks like the brain?



McCulloch & Pitts (1943):
"A logical calculus of the
ideas immanent in nervous
activity"

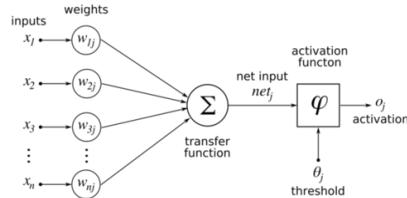
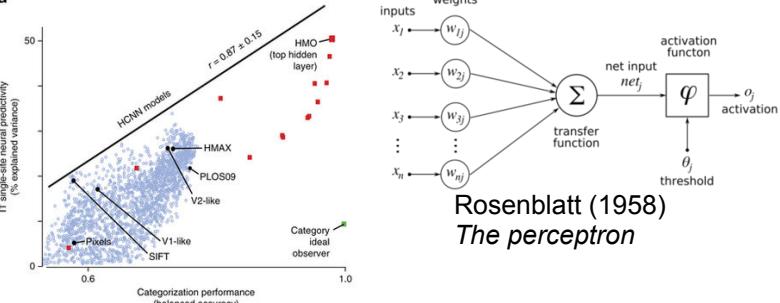


Rosenblatt (1958)
The perceptron



Background: are artificial neural networks like the brain?

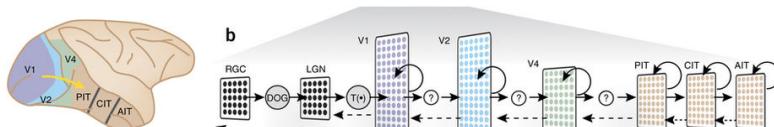
a



Rosenblatt (1958)
The perceptron

a

Stimulus \longrightarrow Encoding \longrightarrow Neurons \longrightarrow Decoding \longrightarrow Behavior

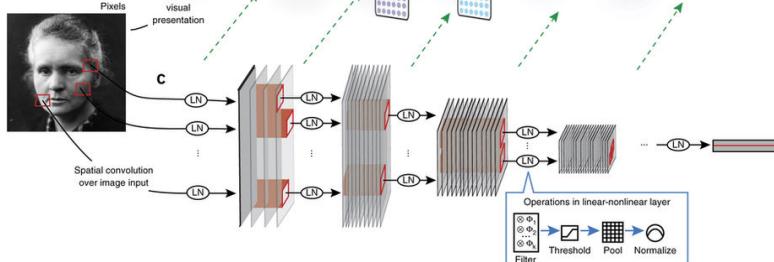


Pixels

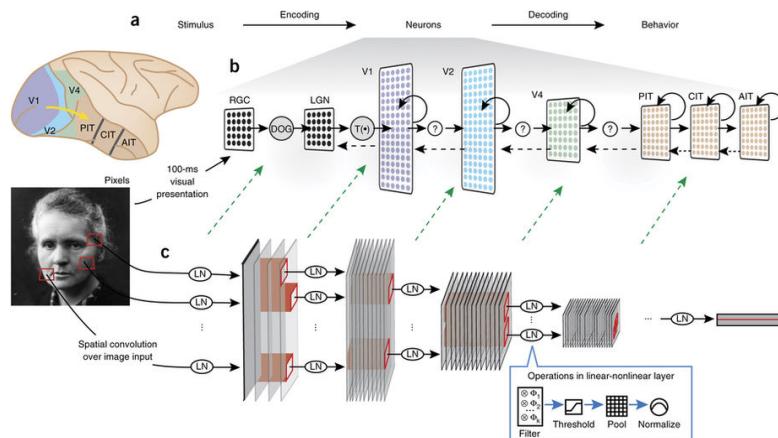
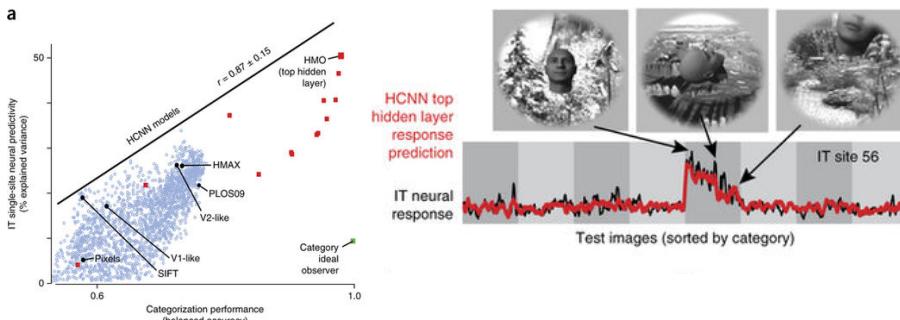
b

c

100-ms visual presentation

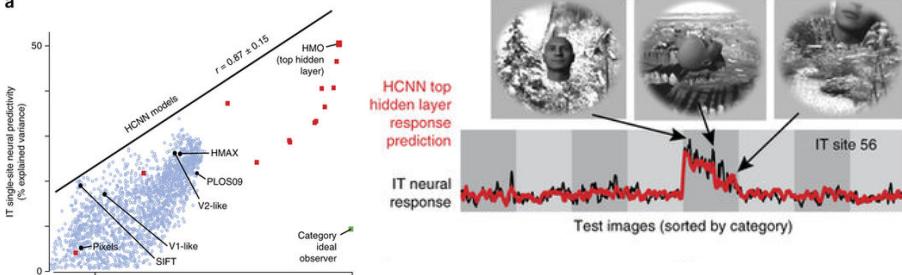


Background: are artificial neural networks like the brain?

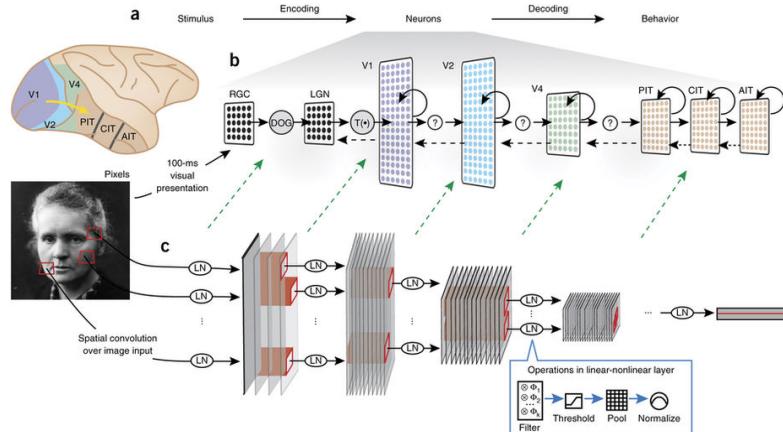


Background: are artificial neural networks like the brain?

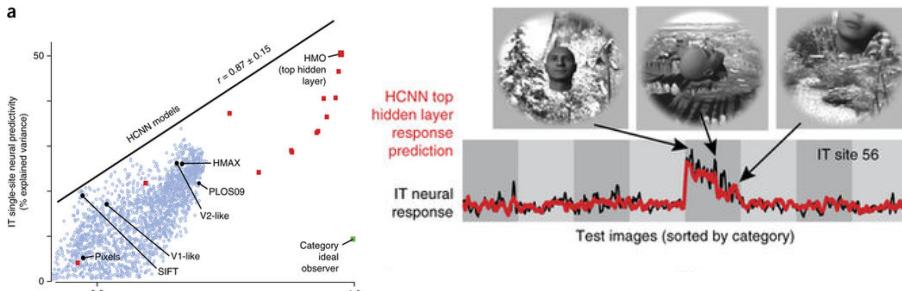
a



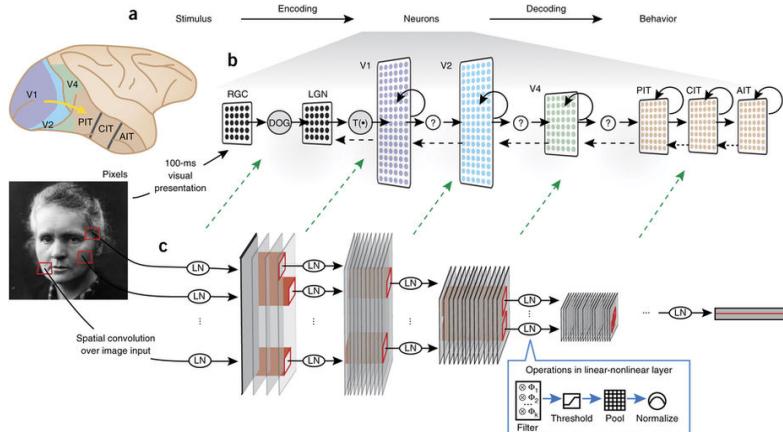
This seems to closely resemble the brain.
Does it work the same way?



Background: are artificial neural networks like the brain?



This seems to closely resemble the brain.
Does it work the same way?



NATURE VOL. 337 12 JANUARY 1989

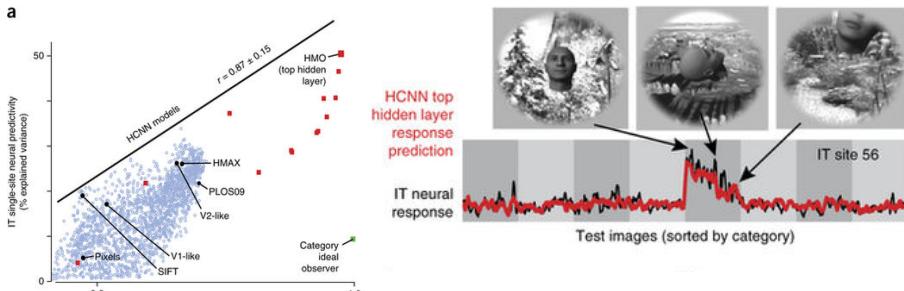
COMMENTARY

129

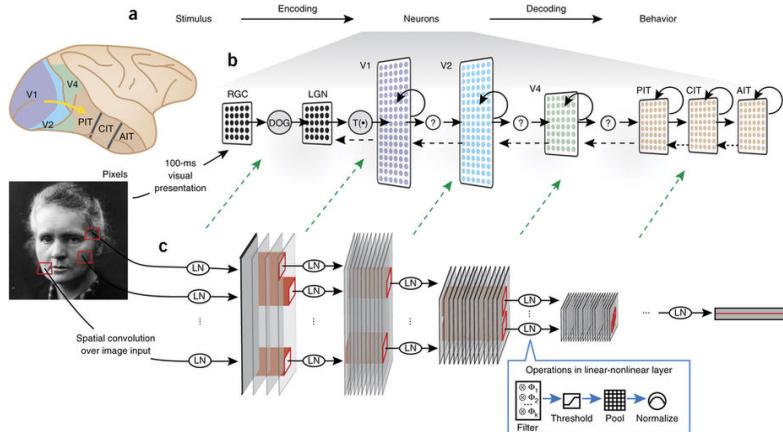
The recent excitement about neural networks

Francis Crick

Background: are artificial neural networks like the brain?



This seems to closely resemble the brain.
Does it work the same way?



NATURE VOL. 337 12 JANUARY 1989

COMMENTARY

129

The recent excitement about neural networks

Francis Crick



François Chollet

Suivre

"Neural networks" are a sad misnomer. They're neither neural nor even networks. They're chains of differentiable, parameterized geometric functions, trained with gradient descent (with gradients obtained via the chain rule). A small set of highschool-level ideas put together

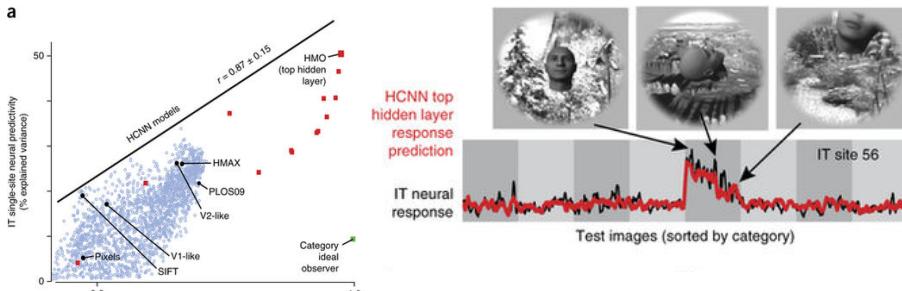
À l'origine en anglais

19:58 - 12 janv. 2018

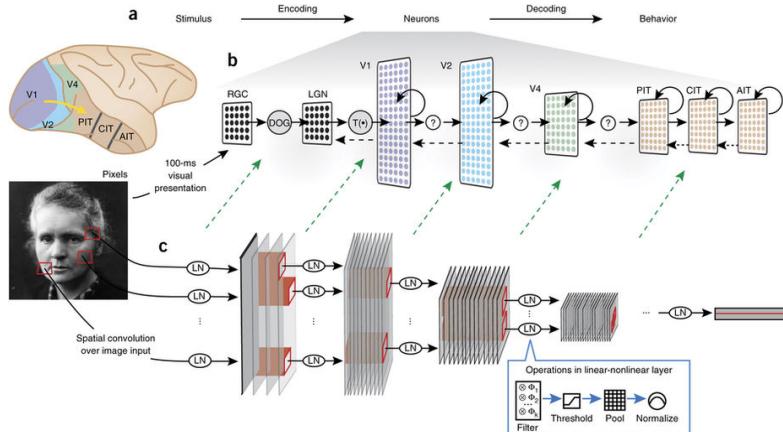
1 314 Retweets 3 388 J'aime

114 1,3 k 3,4 k

Background: are artificial neural networks like the brain?



This seems to closely resemble the brain.
Does it work the same way?



Hassabis et al. (2017); Hinton (2016); Bengio et al. (2016); Yamins & DiCarlo (2016)

NATURE VOL. 337 12 JANUARY 1989

COMMENTARY

129

The recent excitement about neural networks

Francis Crick



François Chollet

Suivre

"Neural networks" are a sad misnomer. They're neither neural nor even networks. They're chains of differentiable, parameterized geometric functions, trained with gradient descent (with gradients obtained via the chain rule). A small set of highschool-level ideas put together

À l'origine en anglais

19:58 - 12 janv. 2018

1 314 Retweets 3 388 J'aime

114 1,3 k 3,4 k

Central argument: Backpropagation is biologically unrealistic.

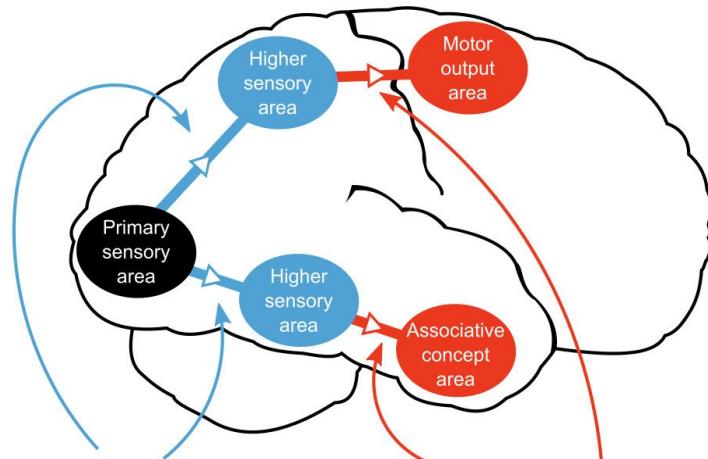
Crick (1987); Harris (2008); Urbanczik & Senn (2009)

Outline

- Introduction: comparing the brain to a neural network
- **The credit assignment problem in learning**
- Primer on the machine learning solution: back-propagation
- Biology's solution? Deep learning with segregated dendrites
- Results
- Discussion

The credit-assignment problem

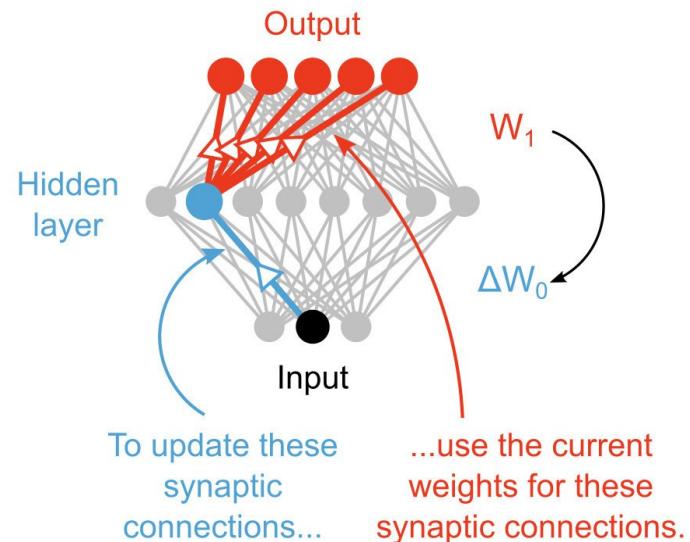
A The "credit assignment" problem



The behavioral effects
of changes to these
synaptic connections...

...depend on the
status of these
synaptic connections.

B The backpropagation solution
(AKA "weight transport")

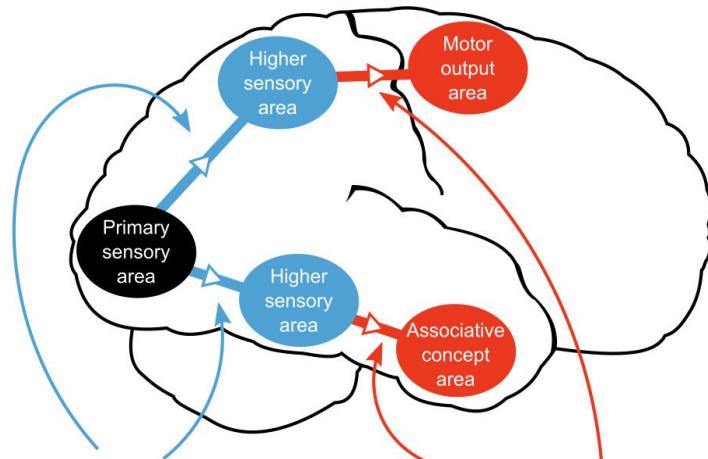


To update these
synaptic
connections...

...use the current
weights for these
synaptic connections.

The credit-assignment problem

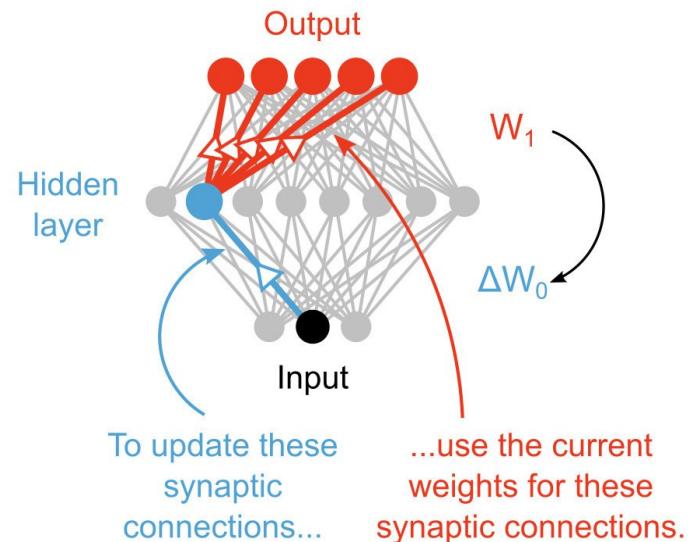
A The "credit assignment" problem



The behavioral effects
of changes to these
synaptic connections...

...depend on the
status of these
synaptic connections.

B The backpropagation solution
(AKA "weight transport")



To update these
synaptic
connections...

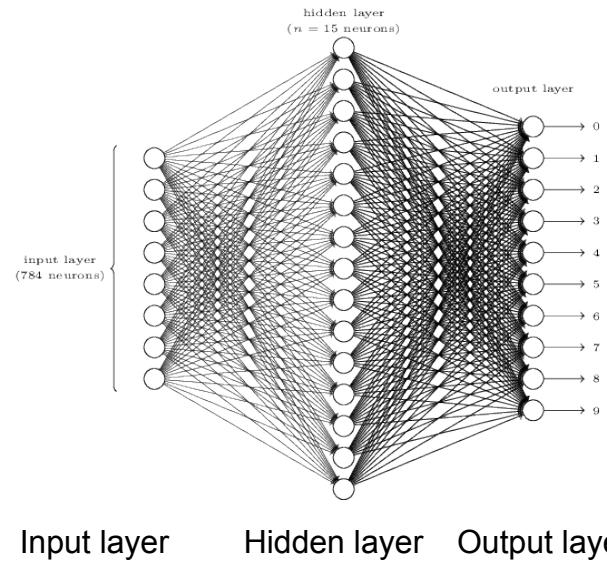
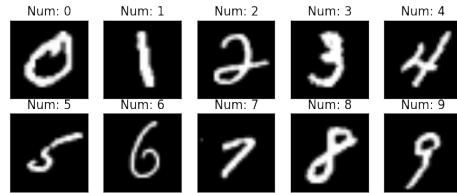
...use the current
weights for these
synaptic connections.

Unrealistic in real brains!

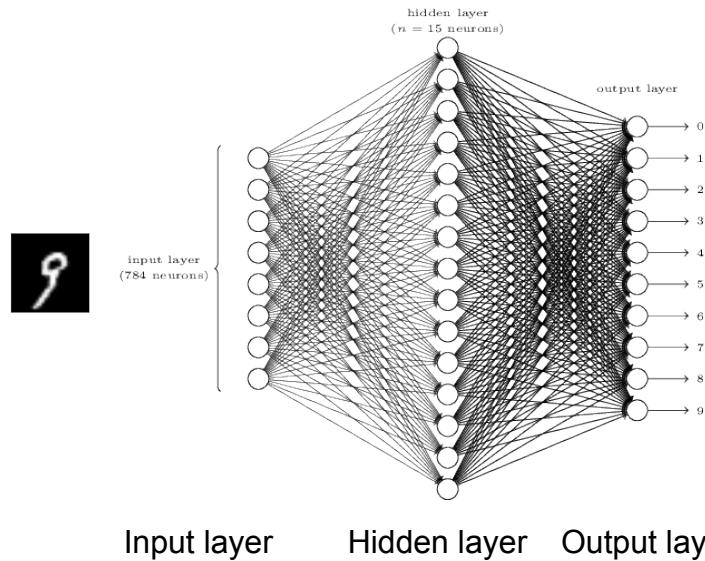
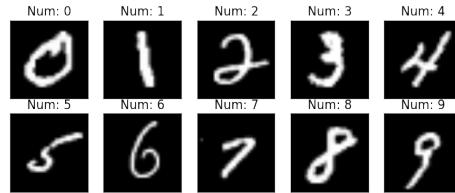
Outline

- Introduction: comparing the brain to a neural network
- The credit assignment problem in learning
- **Primer on the machine learning solution: back-propagation**
- Biology's solution? Deep learning with segregated dendrites
- Results
- Discussion

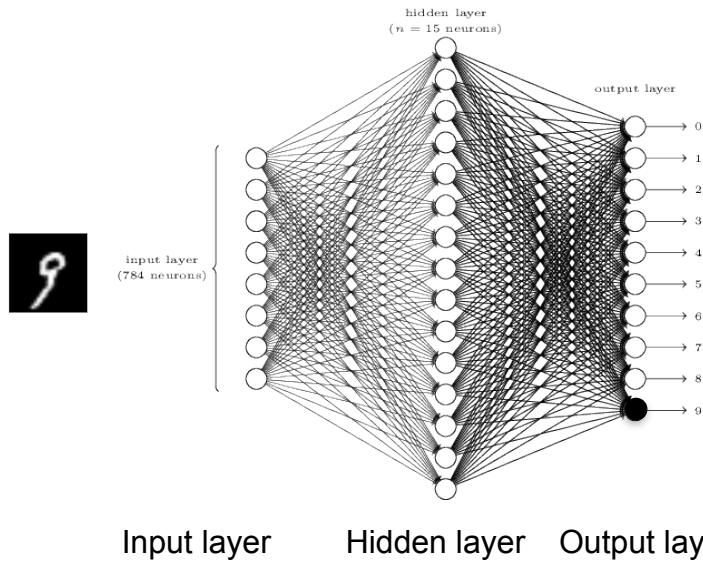
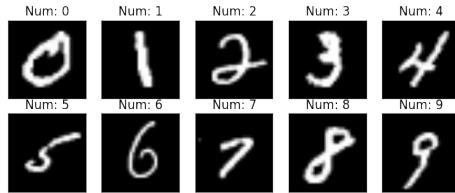
The architecture and terminology of a neural network



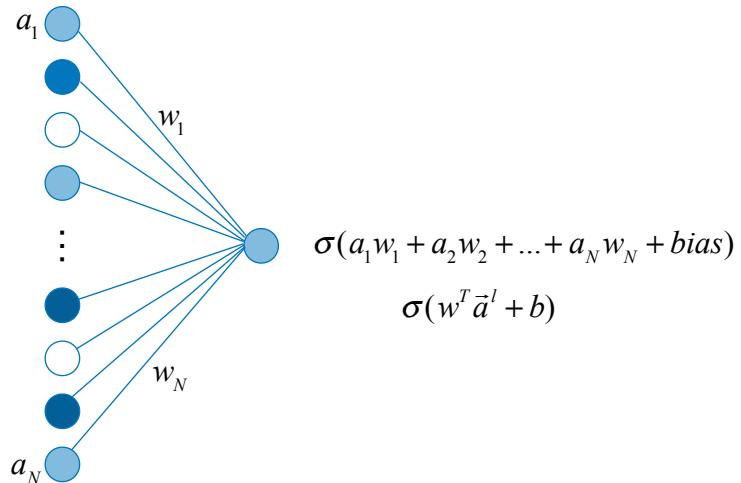
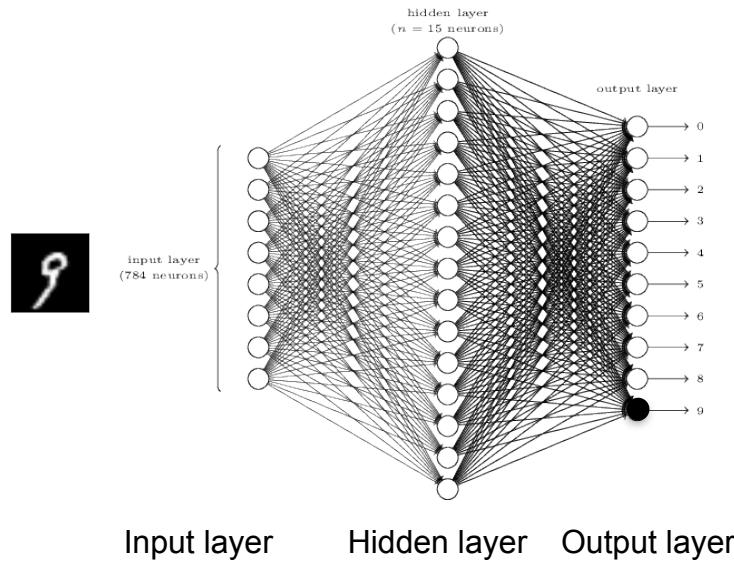
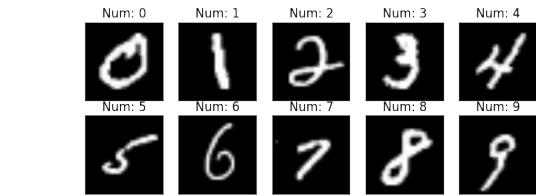
The architecture and terminology of a neural network



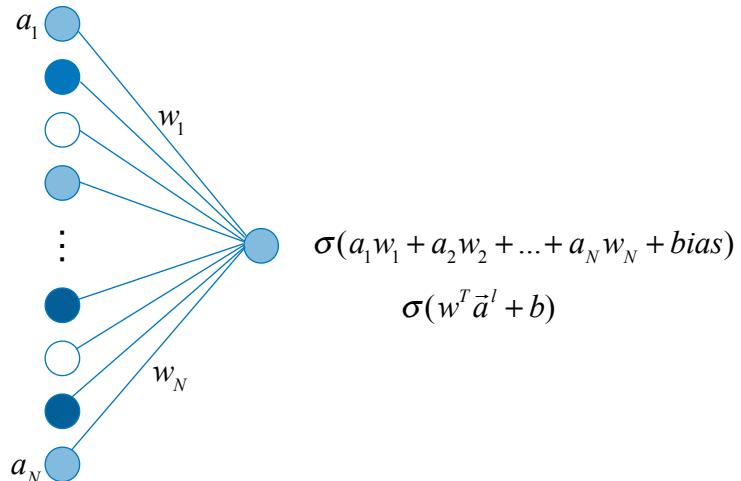
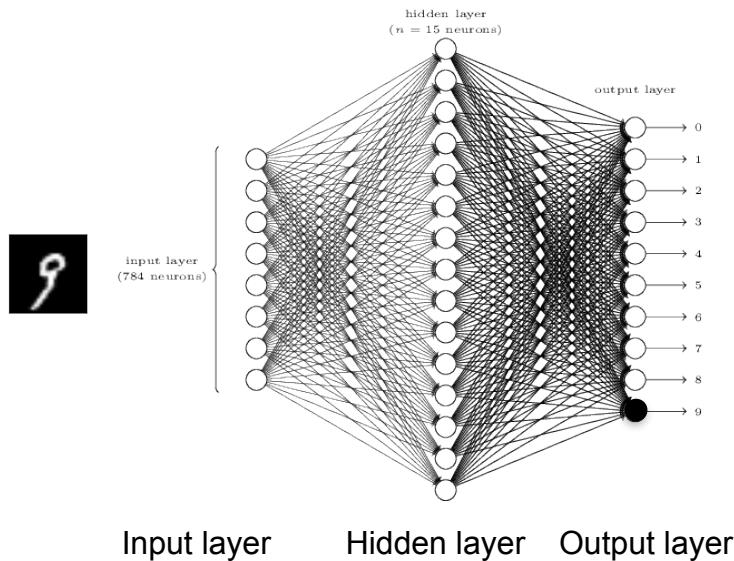
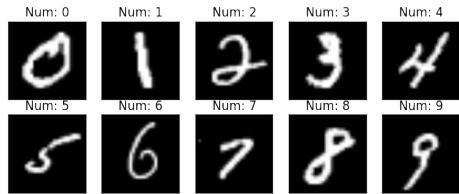
The architecture and terminology of a neural network



The architecture and terminology of a neural network



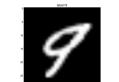
The architecture and terminology of a neural network



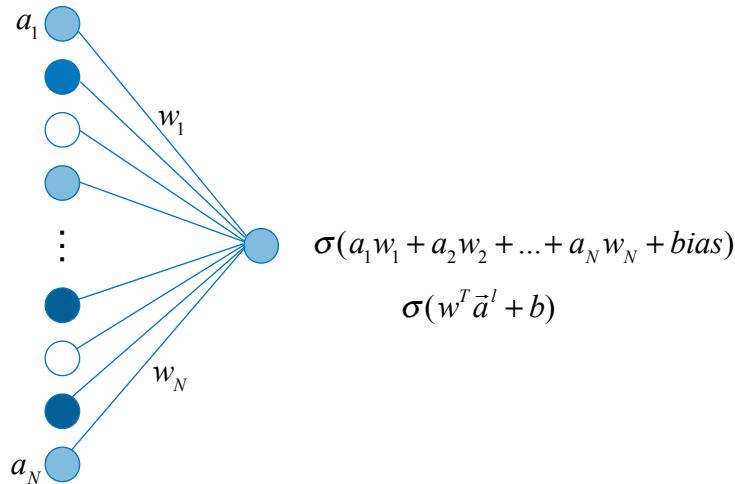
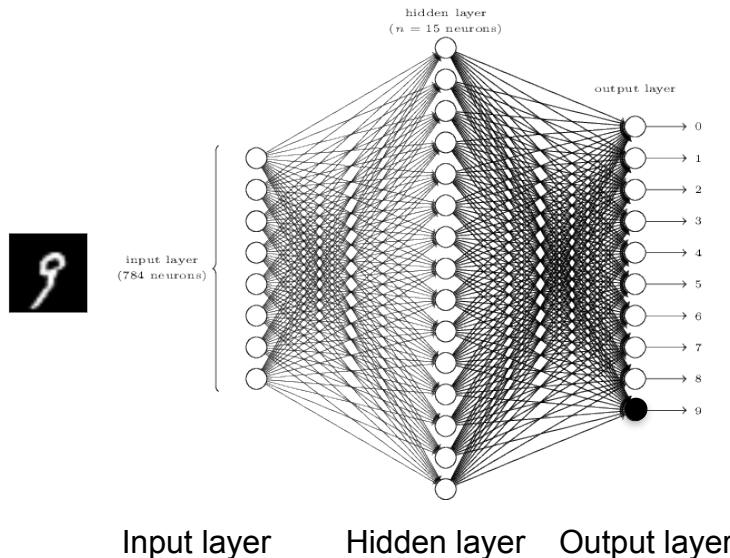
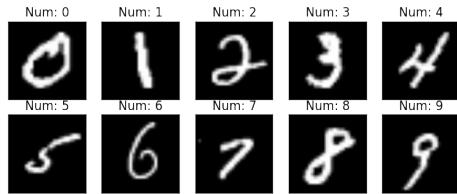
Training

(, 9)

Testing

 ?

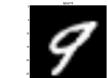
The architecture and terminology of a neural network



Training

(, 9)

Testing

 ?

23,860 parameters for a network with a single hidden layer with 30 neurons! How do we adjust these?

Backpropagation: credit (or blame) assignment in feedforward neural networks

Minimising a cost function

Output
Layer Target y

0	○
1	○
2	○
3	○
4	○
5	○
6	○
7	○
8	○
9	●

9

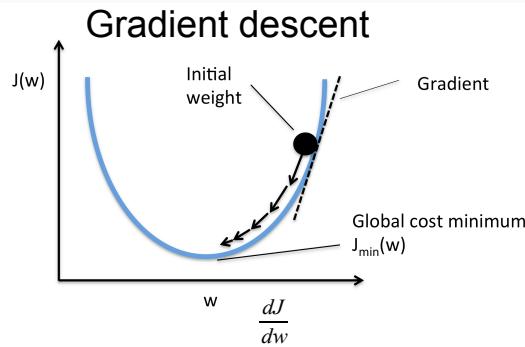


$$J = \sum (a^{(L)} - \bar{y})^2$$

Backpropagation: credit (or blame) assignment in feedforward neural networks

Minimising a cost function

Output Layer	Target y
0	○
1	○
2	○
3	○
4	○
5	○
6	○
7	○
8	○
9	●



9

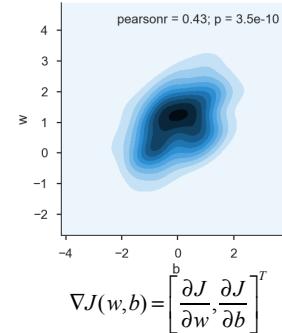
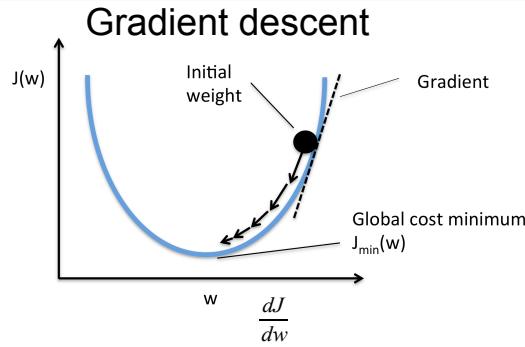


$$J = \sum (a^{(L)} - \bar{y})^2$$

Backpropagation: credit (or blame) assignment in feedforward neural networks

Minimising a cost function

Output Layer	Target y
0	○
1	○
2	○
3	○
4	○
5	○
6	○
7	○
8	○
9	●

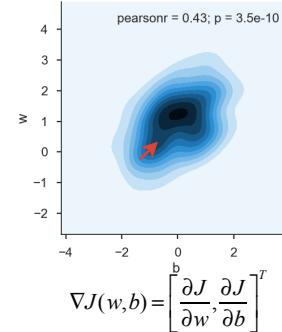
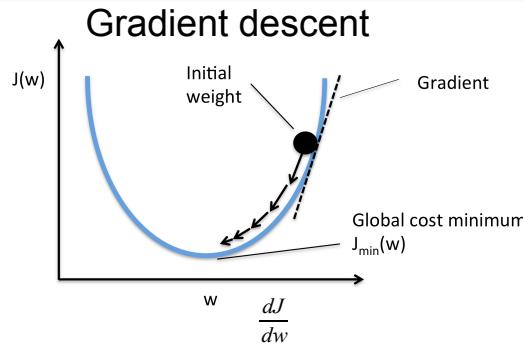


Backpropagation: credit (or blame) assignment in feedforward neural networks

Minimising a cost function

Output Layer	Target y
0	○
1	○
2	○
3	○
4	○
5	○
6	○
7	○
8	○
9	●

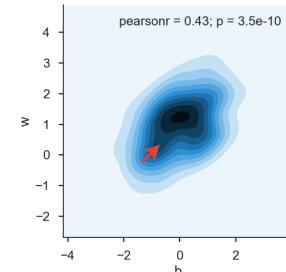
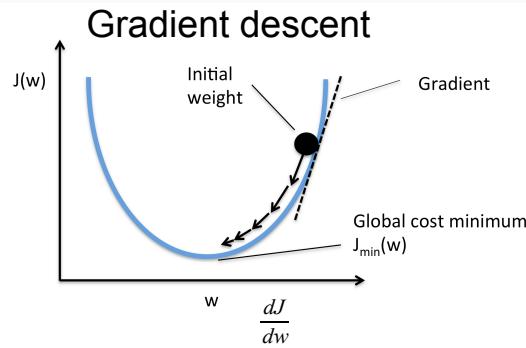
↙



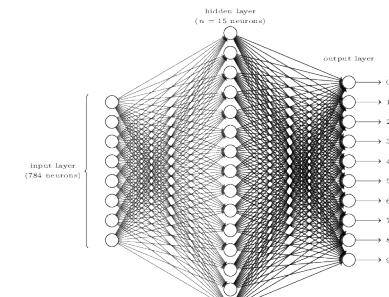
Backpropagation: credit (or blame) assignment in feedforward neural networks

Minimising a cost function

Output Layer	Target y
0	○
1	○
2	○
3	○
4	○
5	○
6	○
7	○
8	○
9	●



$$\nabla J(w, b) = \left[\frac{\partial J}{\partial w}, \frac{\partial J}{\partial b} \right]^T$$



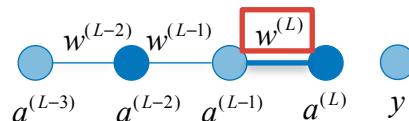
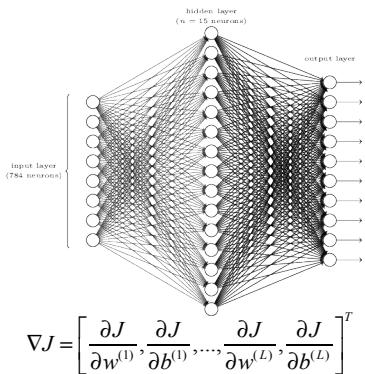
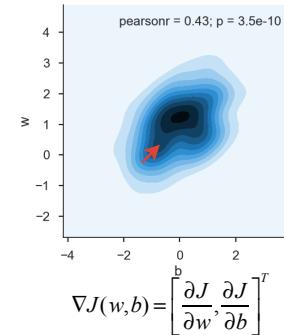
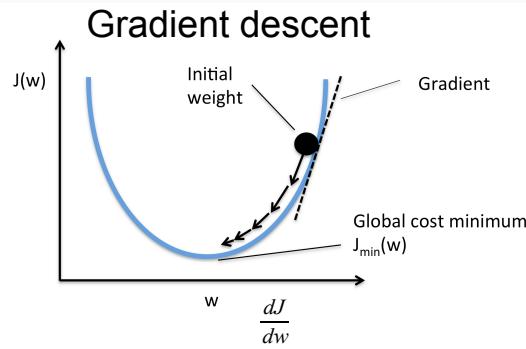
$$\nabla J = \left[\frac{\partial J}{\partial w^{(1)}}, \frac{\partial J}{\partial b^{(1)}}, \dots, \frac{\partial J}{\partial w^{(L)}}, \frac{\partial J}{\partial b^{(L)}} \right]^T$$

$$J = \sum (a^{(L)} - \bar{y})^2$$

Backpropagation: credit (or blame) assignment in feedforward neural networks

Minimising a cost function

Output Layer	Target y
0	○
1	○
2	○
3	○
4	○
5	○
6	○
7	○
8	○
9	●

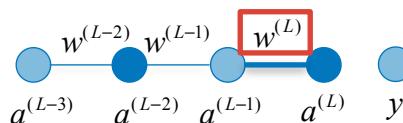
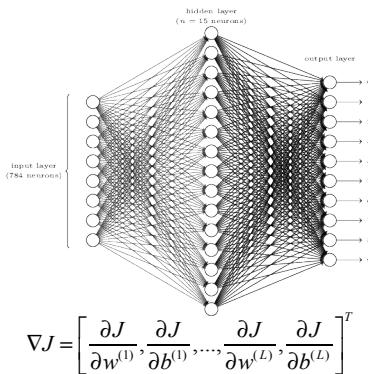
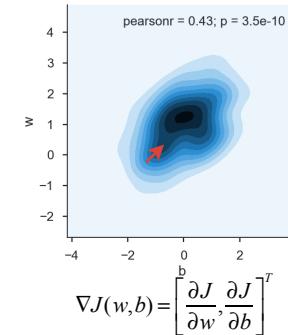
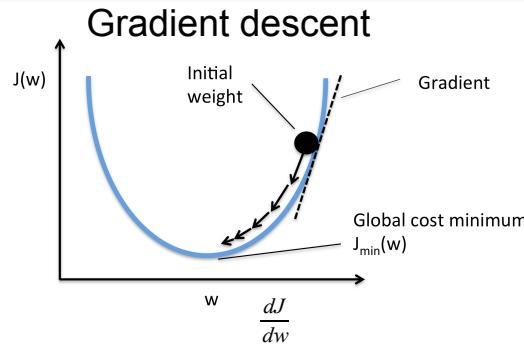


$$J = \sum (a^{(L)} - \bar{y})^2$$

Backpropagation: credit (or blame) assignment in feedforward neural networks

Minimising a cost function

Output Layer	Target y
0	○
1	○
2	○
3	○
4	○
5	○
6	○
7	○
8	○
9	●



How does the cost depend on changes in a specific weight?

$$J = \sum (a^{(L)} - \bar{y})^2$$

Backpropagation: credit (or blame) assignment in feedforward neural networks

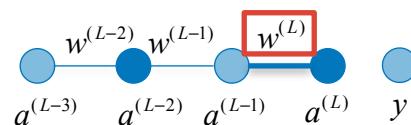
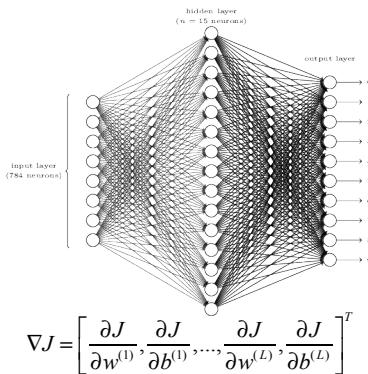
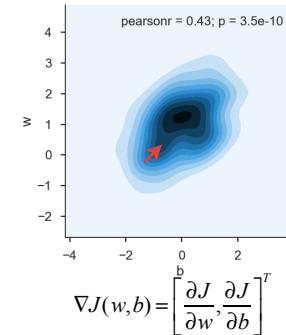
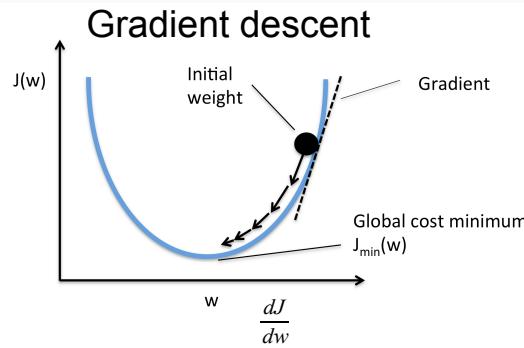
Minimising a cost function

Output Layer Target y

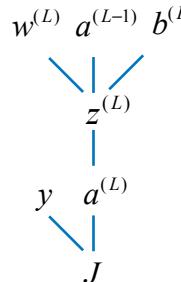
- 0
- 1
- 2
- 3
- 4
- 5
- 6
- 7
- 8
- 9



$$J = \sum (a^{(L)} - \bar{y})^2$$



How does the cost depend on changes in a specific weight?



Backpropagation: credit (or blame) assignment in feedforward neural networks

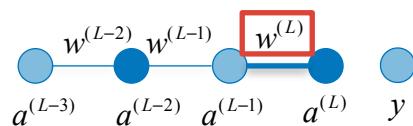
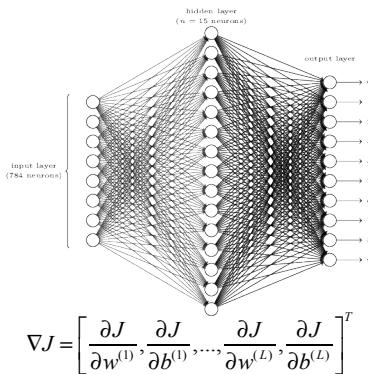
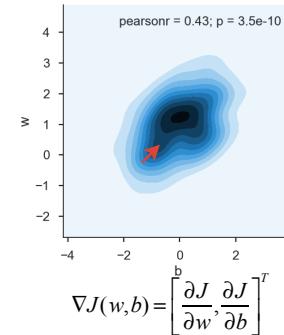
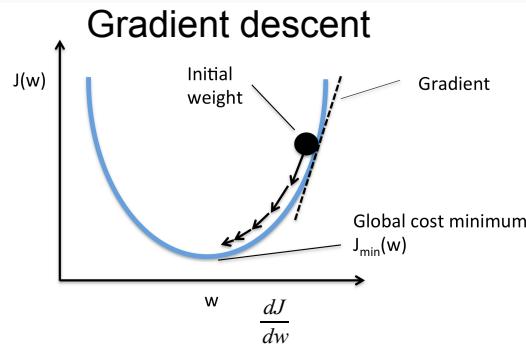
Minimising a cost function

Output Layer Target y

- 0
- 1
- 2
- 3
- 4
- 5
- 6
- 7
- 8
- 9



$$J = \sum (a^{(L)} - \bar{y})^2$$



How does the cost depend on changes in a specific weight?

$$\frac{\partial J}{\partial w^{(L)}} =$$

$w^{(L)} \quad a^{(L-1)} \quad b^{(L)}$

$z^{(L)}$

$y \quad a^{(L)}$

J

Backpropagation: credit (or blame) assignment in feedforward neural networks

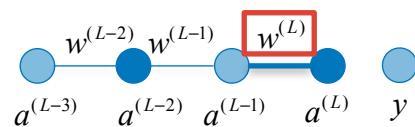
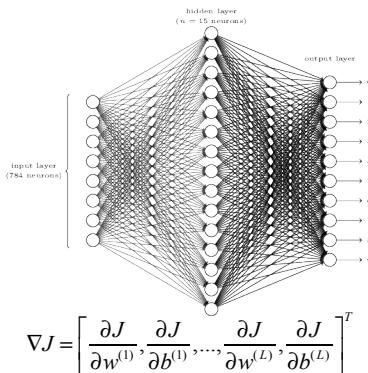
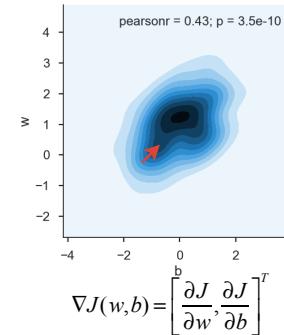
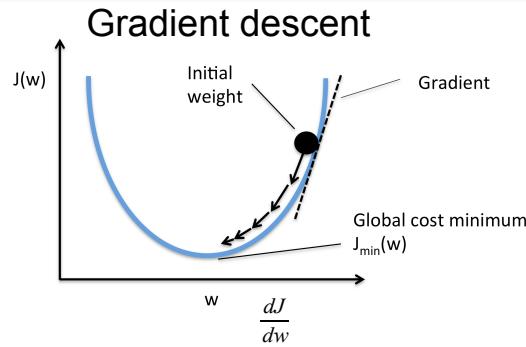
Minimising a cost function

Output Layer Target y

- 0
- 1
- 2
- 3
- 4
- 5
- 6
- 7
- 8
- 9



$$J = \sum (a^{(L)} - \bar{y})^2$$



How does the cost depend on changes in a specific weight?

$$\frac{\partial J}{\partial w^{(L)}} = \frac{\partial z^{(L)}}{\partial w^{(L)}}$$

$w^{(L)}$ $a^{(L-1)}$ $b^{(L)}$

$z^{(L)}$

y $a^{(L)}$

J

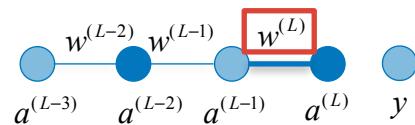
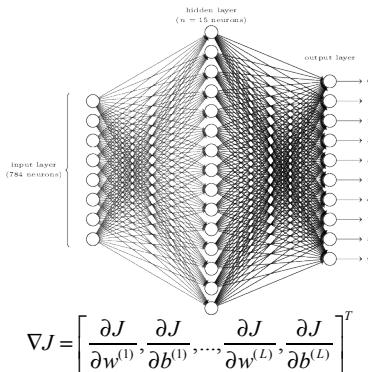
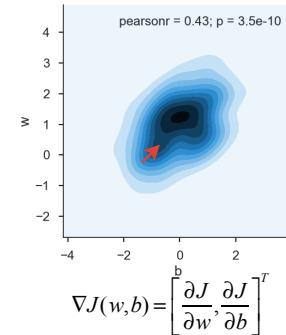
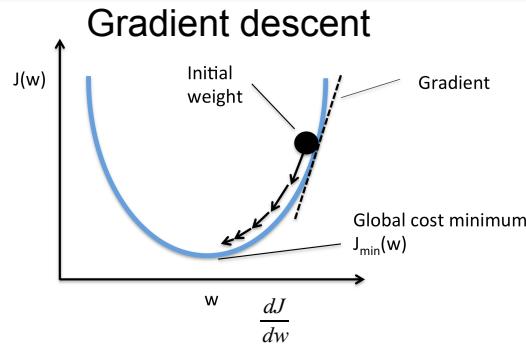
Backpropagation: credit (or blame) assignment in feedforward neural networks

Minimising a cost function

Output Layer	Target y
0	○
1	○
2	○
3	○
4	○
5	○
6	○
7	○
8	○
9	●

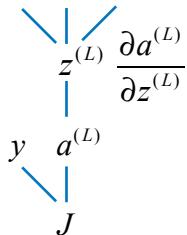


$$J = \sum (a^{(L)} - \bar{y})^2$$



How does the cost depend on changes in a specific weight?

$$\frac{\partial J}{\partial w^{(L)}} = \frac{\partial z^{(L)}}{\partial w^{(L)}}$$



Backpropagation: credit (or blame) assignment in feedforward neural networks

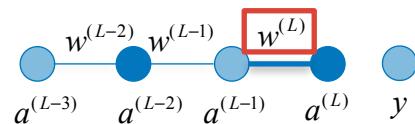
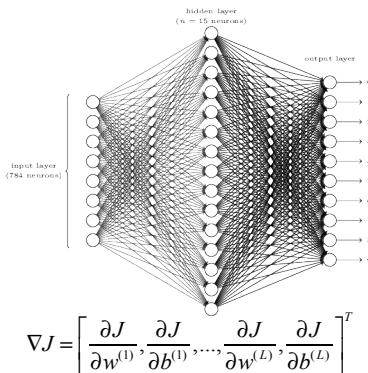
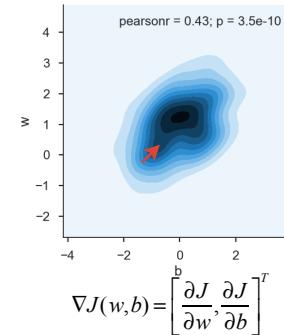
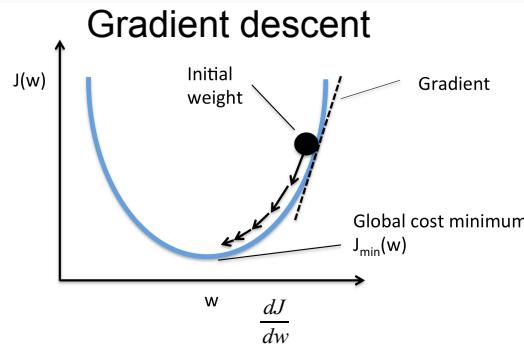
Minimising a cost function

Output Layer Target y

- 0
- 1
- 2
- 3
- 4
- 5
- 6
- 7
- 8
- 9



$$J = \sum (a^{(L)} - \bar{y})^2$$



How does the cost depend on changes in a specific weight?

$$\frac{\partial J}{\partial w^{(L)}} = \frac{\partial z^{(L)}}{\partial w^{(L)}} \frac{\partial a^{(L)}}{\partial z^{(L)}}$$

$w^{(L)}$ $a^{(L-1)}$ $b^{(L)}$

$z^{(L)}$

y $a^{(L)}$

J

Backpropagation: credit (or blame) assignment in feedforward neural networks

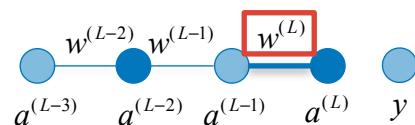
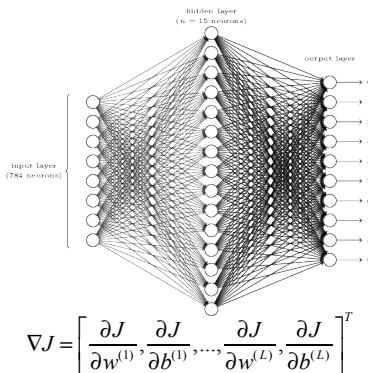
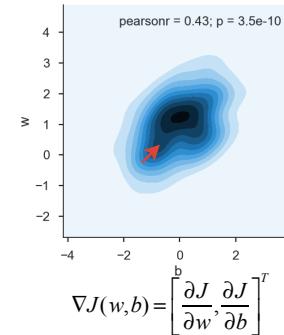
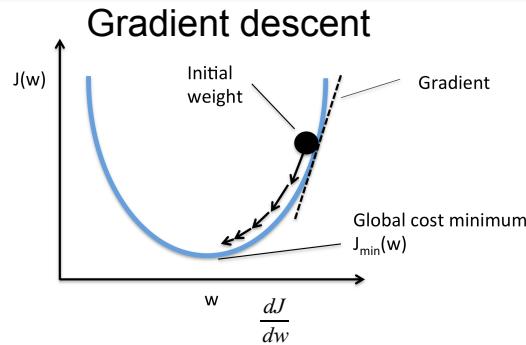
Minimising a cost function

Output Layer Target y

- 0
- 1
- 2
- 3
- 4
- 5
- 6
- 7
- 8
- 9



$$J = \sum (a^{(L)} - \bar{y})^2$$



How does the cost depend on changes in a specific weight?

$$\frac{\partial J}{\partial w^{(L)}} = \frac{\partial z^{(L)}}{\partial w^{(L)}} \frac{\partial a^{(L)}}{\partial z^{(L)}}$$

$w^{(L)}$ $a^{(L-1)}$ $b^{(L)}$

$z^{(L)}$

y $a^{(L)}$ $\frac{\partial J}{\partial a^{(L)}}$

Backpropagation: credit (or blame) assignment in feedforward neural networks

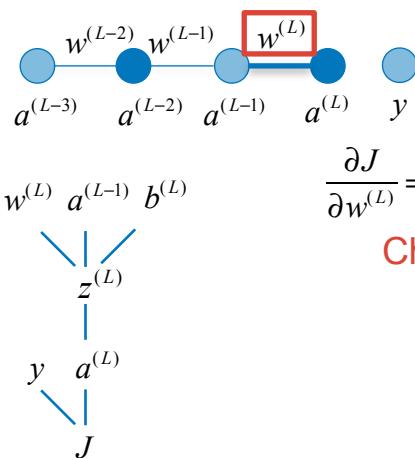
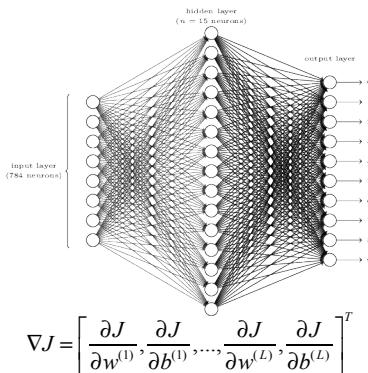
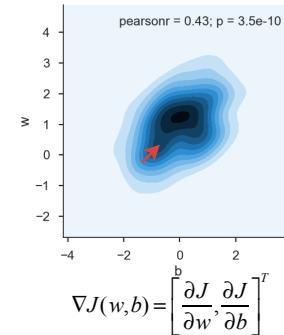
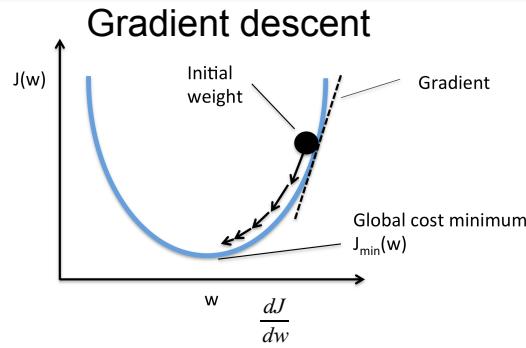
Minimising a cost function

Output Layer Target y

- 0
- 1
- 2
- 3
- 4
- 5
- 6
- 7
- 8
- 9



$$J = \sum (a^{(L)} - \bar{y})^2$$



How does the cost depend on changes in a specific weight?

$$\frac{\partial J}{\partial w^{(L)}} = \frac{\partial z^{(L)}}{\partial w^{(L)}} \frac{\partial a^{(L)}}{\partial z^{(L)}} \frac{\partial J}{\partial a^{(L)}}$$

Chain rule!

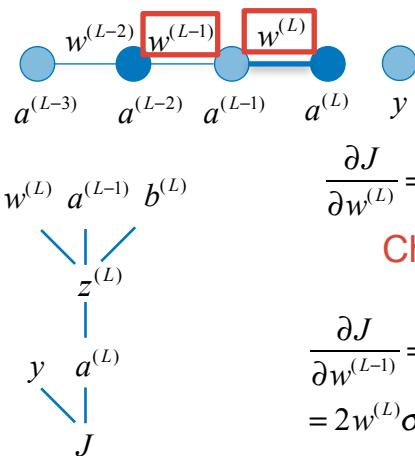
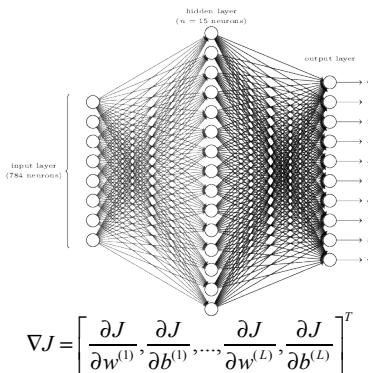
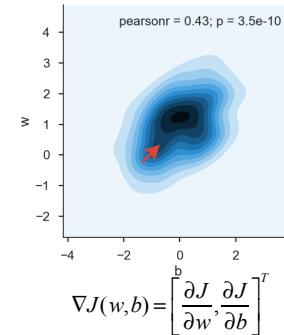
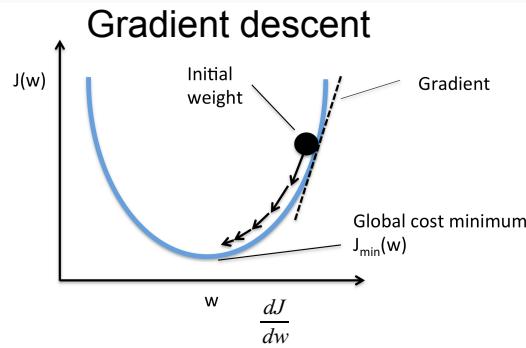
Backpropagation: credit (or blame) assignment in feedforward neural networks

Minimising a cost function

Output Layer	Target y
0	○
1	○
2	○
3	○
4	○
5	○
6	○
7	○
8	○
9	●



$$J = \sum (a^{(L)} - \bar{y})^2$$



How does the cost depend on changes in a specific weight?

$$\frac{\partial J}{\partial w^{(L)}} = \frac{\partial z^{(L)}}{\partial w^{(L)}} \frac{\partial a^{(L)}}{\partial z^{(L)}} \frac{\partial J}{\partial a^{(L)}}$$

Chain rule!

$$\begin{aligned} \frac{\partial J}{\partial w^{(L-1)}} &= \frac{\partial z^{(L-1)}}{\partial w^{(L-1)}} \frac{\partial a^{(L-1)}}{\partial z^{(L-1)}} \frac{\partial J}{\partial a^{(L-1)}} \\ &= 2w^{(L)}\sigma'(z^{(L)})(a^{(L)} - y) \end{aligned}$$

Backpropagation: credit (or blame) assignment in feedforward neural networks

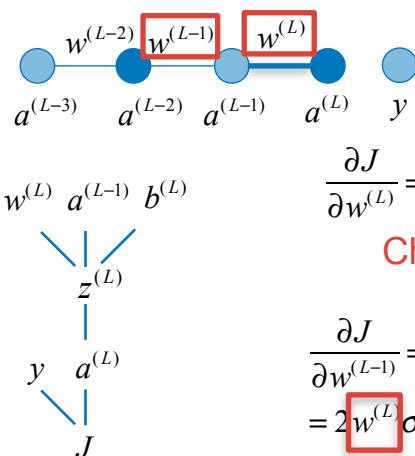
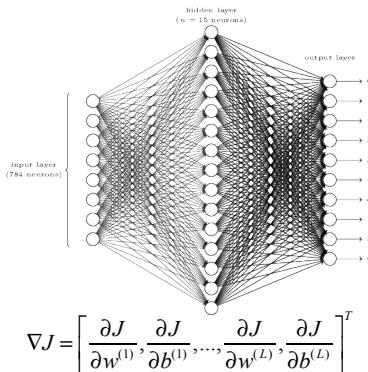
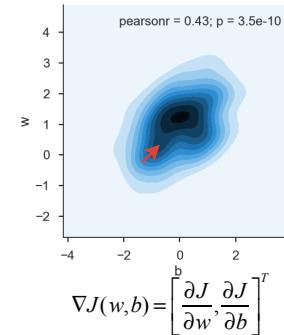
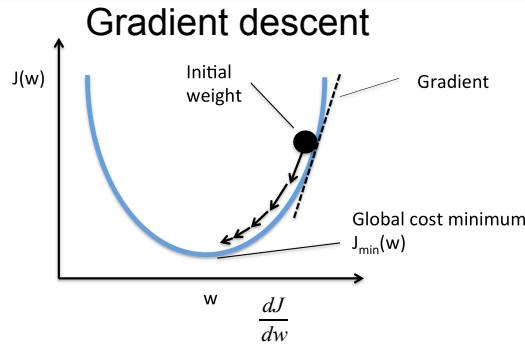
Minimising a cost function

Output Layer Target y

	0	
	1	
	2	
	3	
	4	
	5	
	6	
	7	
	8	
	9	



$$J = \sum (a^{(L)} - \bar{y})^2$$



How does the cost depend on changes in a specific weight?

$$\frac{\partial J}{\partial w^{(L)}} = \frac{\partial z^{(L)}}{\partial w^{(L)}} \frac{\partial a^{(L)}}{\partial z^{(L)}} \frac{\partial J}{\partial a^{(L)}}$$

Chain rule!

$$\begin{aligned} \frac{\partial J}{\partial w^{(L-1)}} &= \frac{\partial z^{(L-1)}}{\partial w^{(L-1)}} \frac{\partial a^{(L-1)}}{\partial z^{(L-1)}} \frac{\partial J}{\partial a^{(L-1)}} \\ &= 2w^{(L)} \sigma'(z^{(L)}) (a^{(L)} - y) \end{aligned}$$

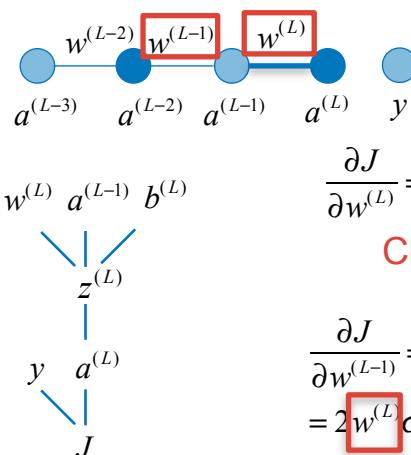
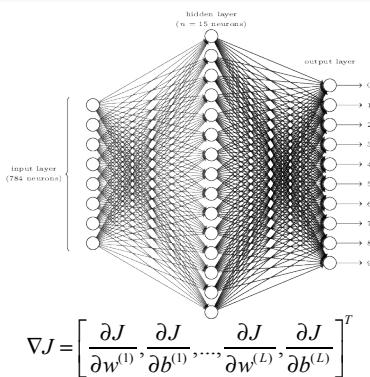
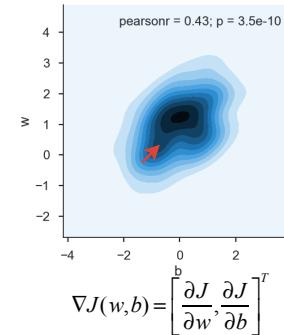
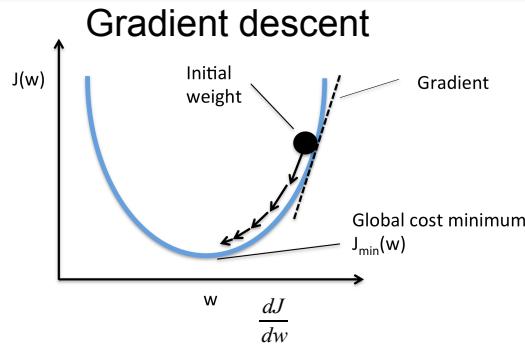
Backpropagation: credit (or blame) assignment in feedforward neural networks

Minimising a cost function

Output Layer Target y

- 0
- 1
- 2
- 3
- 4
- 5
- 6
- 7
- 8
- 9

$$J = \sum (a^{(L)} - \bar{y})^2$$



How does the cost depend on changes in a specific weight?

$$\frac{\partial J}{\partial w^{(L)}} = \frac{\partial z^{(L)}}{\partial w^{(L)}} \frac{\partial a^{(L)}}{\partial z^{(L)}} \frac{\partial J}{\partial a^{(L)}}$$

Chain rule!

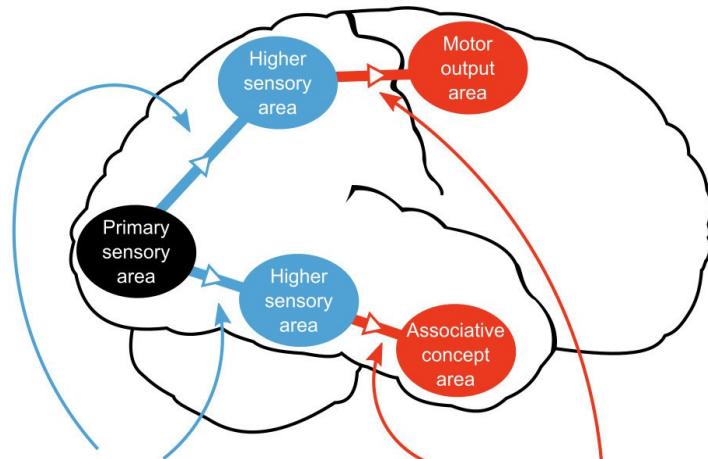
$$\begin{aligned} \frac{\partial J}{\partial w^{(L-1)}} &= \frac{\partial z^{(L-1)}}{\partial w^{(L-1)}} \frac{\partial a^{(L-1)}}{\partial z^{(L-1)}} \frac{\partial J}{\partial a^{(L-1)}} \\ &= 2w^{(L)} \sigma'(z^{(L)}) (a^{(L)} - y) \end{aligned}$$

Backpropagation summary:

- Define a cost function
- Initialise weights and biases randomly
- Find gradient using chain rule, working backwards through the network
- Nudge weights and biases in the opposite direction
- Repeat for all training examples
- **Change in shallower layer weights depends on ALL deeper layer weights!**

The credit-assignment problem

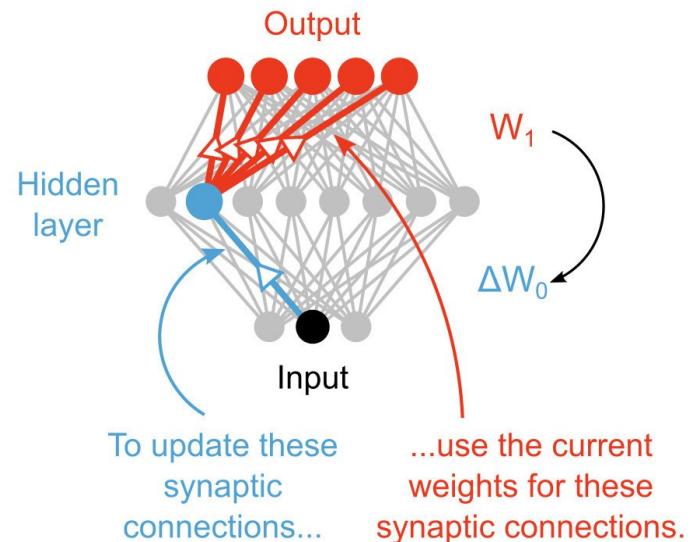
A The "credit assignment" problem



The behavioral effects
of changes to these
synaptic connections...

...depend on the
status of these
synaptic connections.

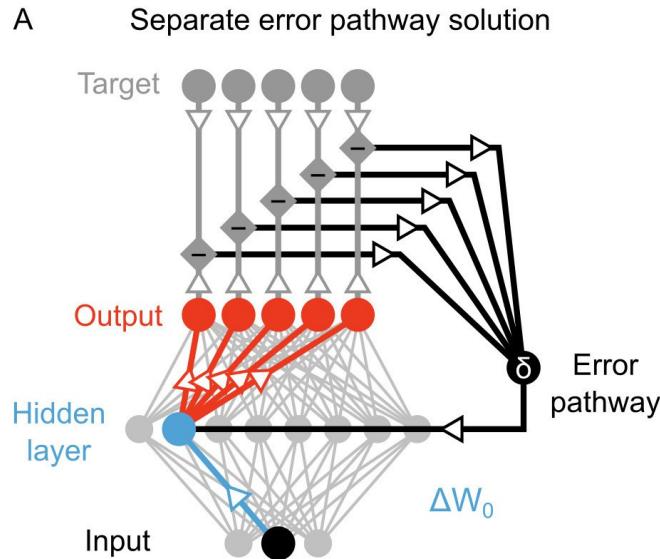
B The backpropagation solution
(AKA "weight transport")



To update these
synaptic
connections...

...use the current
weights for these
synaptic connections.

Backpropagation in the brain with a separate error pathway is problematic

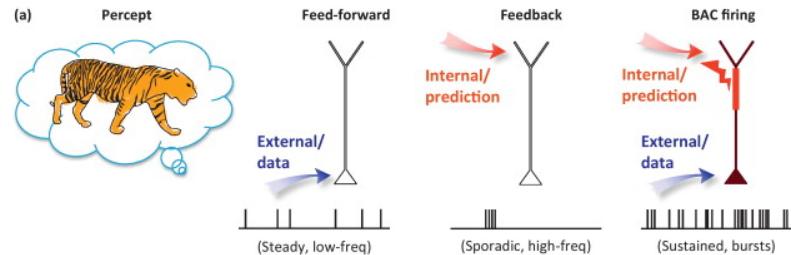
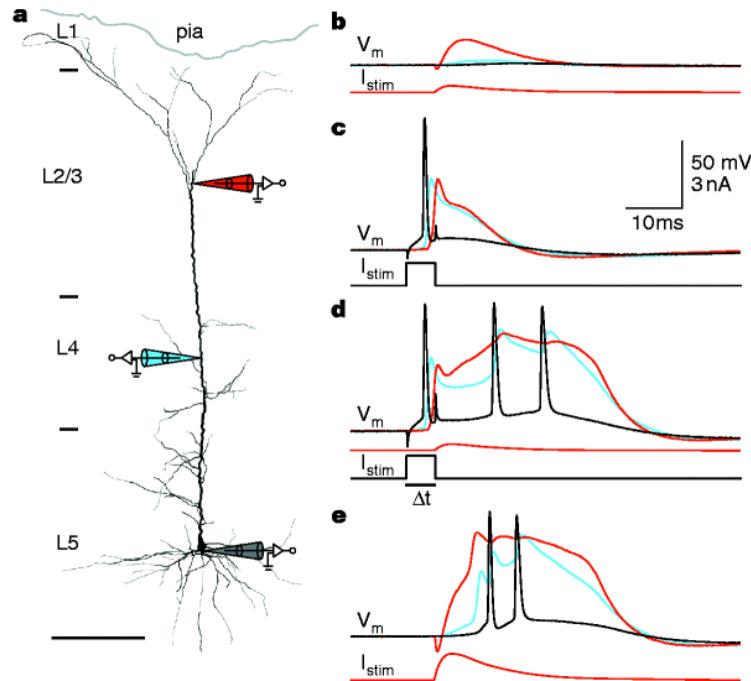


- An error neuron for every hidden layer neuron?
- The error signal is signed. Using a baseline firing rate would lead to noise when the error gets smaller

Outline

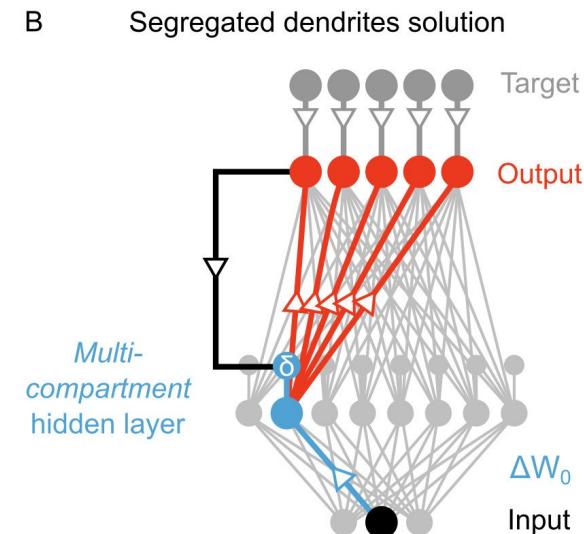
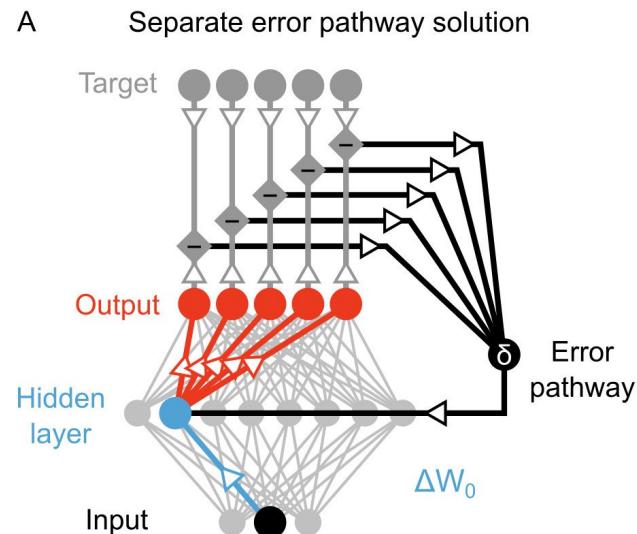
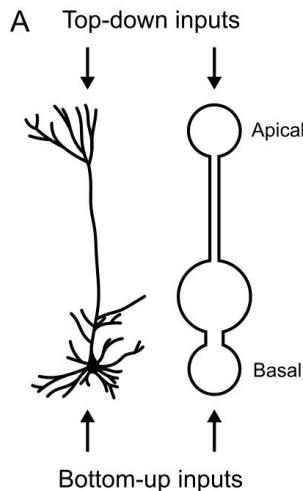
- Introduction: comparing the brain to a neural network
- The credit assignment problem in learning
- Primer on the machine learning solution: back-propagation
- **Biology's solution? Deep learning with segregated dendrites**
- Results
- Discussion

Calcium-driven dendritic spikes for associating inputs to an internal model

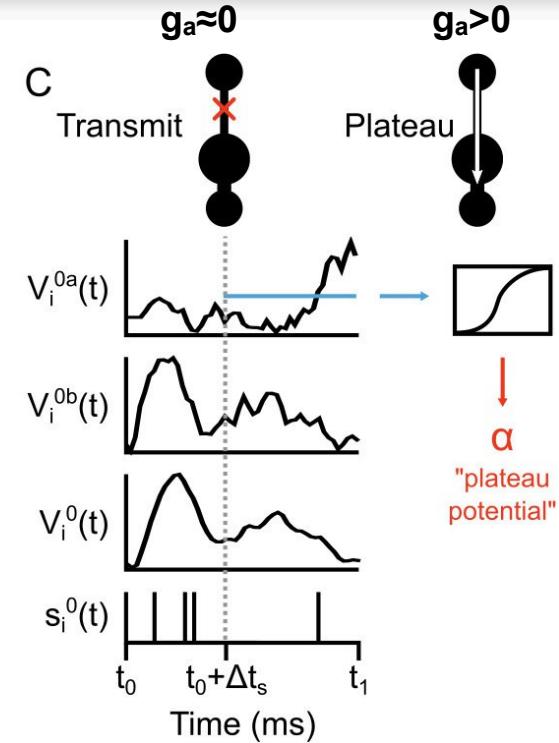
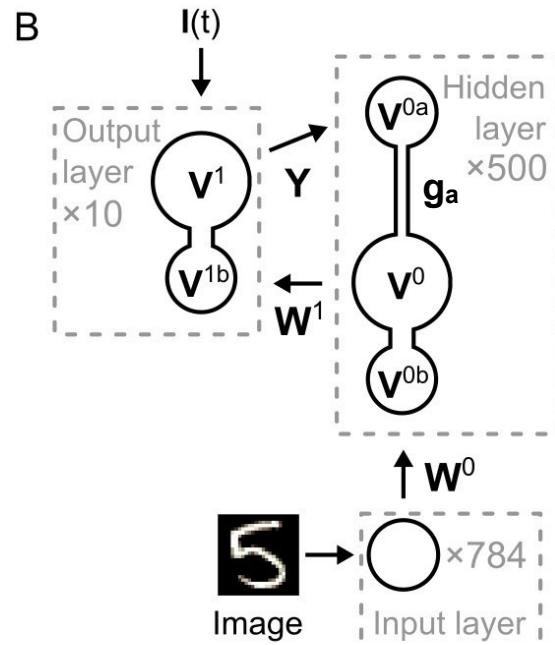
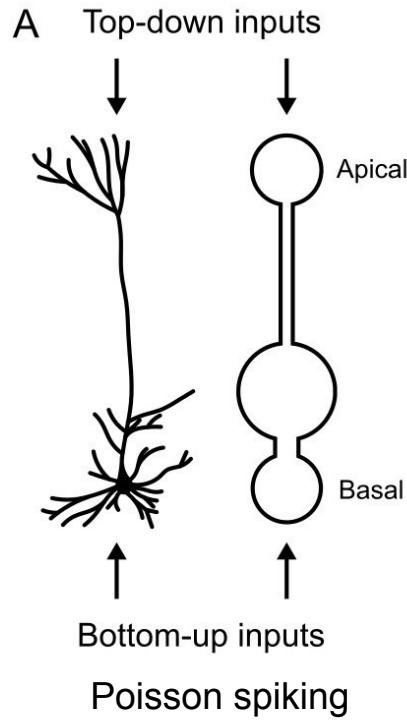


- Real pyramidal neurons have electronically distant apical dendrites
- Apical to soma transduction requires active propagation (through voltage gated Ca^{2+} channels).
- Bottom-up input arrives at basal dendrites, higher order feedback at apical dendrites (to some extent)
- Simultaneous apical and basal stimulation leads to plateau potentials

Solving the credit assignment problem with segregated dendrites



Basic model architecture: neurons with segregated dendritic compartments



$$\tau \frac{dV_i^0(t)}{dt} = -V_i^0(t) + \frac{g_b}{g_l} (V_i^{0b}(t) - V_i^0(t)) + \frac{g_a}{g_l} (V_i^{0a}(t) - V_i^0(t))$$

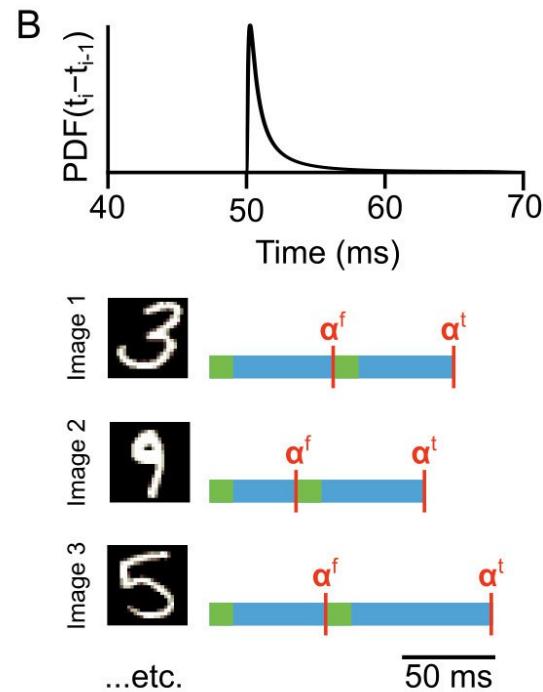
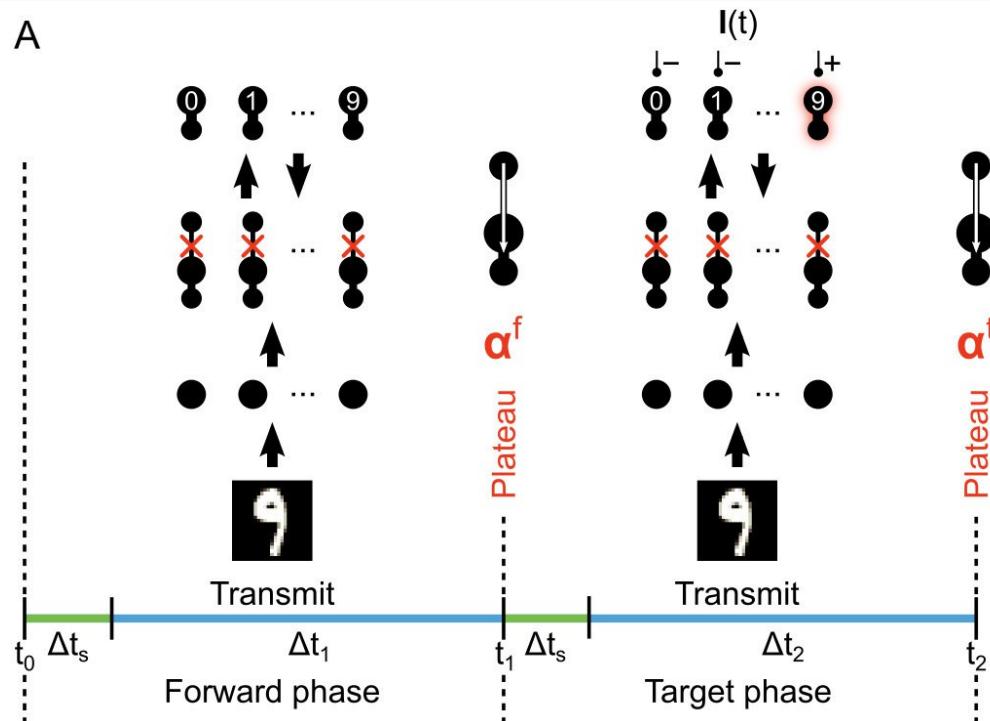
$$V_i^{0b}(t) = \sum_{j=1}^{\ell} W_{ij}^0 s_j^{\text{input}}(t) + b_i^0$$

$$V_i^{0a}(t) = \sum_{j=1}^n Y_{ij} s_j^1(t)$$

$$\tau \frac{dV_i^1(t)}{dt} = -V_i^1(t) + \frac{g_d}{g_l} (V_i^{1b}(t) - V_i^1(t)) + I_i(t)$$

$$V_i^{1b}(t) = \sum_{j=1}^{\ell} W_{ij}^1 s_j^0(t) + b_i^1$$

Network learning in two distinct phases



$$\alpha_i^f = \sigma \left(\frac{1}{\Delta t_1} \int_{t_1 - \Delta t_1}^{t_1} V_i^{0a}(t) dt \right)$$

$$\alpha_i^t = \sigma \left(\frac{1}{\Delta t_2} \int_{t_2 - \Delta t_2}^{t_2} V_i^{0a}(t) dt \right)$$

Defining the local loss functions

*Actual loss functions more complicated and shown in methods

Defining the local loss functions

Output layer target and loss

$$\begin{aligned}\phi_i^{1*} &= \overline{\phi_i^1}^t \\ &= \frac{1}{\Delta t_2} \int_{t_1 + \Delta t_s}^{t_2} \phi_i^1(t) dt\end{aligned}$$

$$\begin{aligned}L^1 &\approx \|\phi^{1*} - \overline{\phi^1}^f\|_2^2 \\ &= \|\overline{\phi^1}^t - \overline{\phi^1}^f\|_2^2 \\ &= \left\| \left| \frac{1}{\Delta t_2} \int_{t_1 + \Delta t_s}^{t_2} \phi^1(t) dt - \frac{1}{\Delta t_1} \int_{t_0 + \Delta t_s}^{t_1} \phi^1(t) dt \right| \right\|_2^2\end{aligned}$$

Defining the local loss functions

Output layer target and loss

$$\begin{aligned}\phi_i^{1*} &= \overline{\phi}_i^t \\ &= \frac{1}{\Delta t_2} \int_{t_1 + \Delta t_s}^{t_2} \phi_i^1(t) dt\end{aligned}$$

$$\begin{aligned}L^1 &\approx \|\phi^{1*} - \overline{\phi}^f\|_2^2 \\ &= \|\overline{\phi}^t - \overline{\phi}^f\|_2^2 \\ &= \left\| \left| \frac{1}{\Delta t_2} \int_{t_1 + \Delta t_s}^{t_2} \phi^1(t) dt - \frac{1}{\Delta t_1} \int_{t_0 + \Delta t_s}^{t_1} \phi^1(t) dt \right| \right\|_2^2\end{aligned}$$

Hidden layer target and loss

$$\phi_i^{0*} = \overline{\phi}_i^f + \alpha_i^t - \alpha_i^f$$

$$\begin{aligned}L^0 &\approx \|\phi^{0*} - \overline{\phi}^0\|_2^2 \\ &= \|\overline{\phi}^f + \alpha^t - \alpha_i^f - \overline{\phi}^0\|_2^2 \\ &= \|\alpha^t - \alpha^f\|_2^2\end{aligned}$$

*Actual loss functions more complicated and shown in methods

Defining the local loss functions

Output layer target and loss

$$\begin{aligned}\phi_i^{1*} &= \overline{\phi}_i^t \\ &= \frac{1}{\Delta t_2} \int_{t_1+\Delta t_s}^{t_2} \phi_i^1(t) dt\end{aligned}$$

$$\begin{aligned}L^1 &\approx \|\phi^{1*} - \overline{\phi}^f\|_2^2 \\ &= \|\overline{\phi}^t - \overline{\phi}^f\|_2^2 \\ &= \left\| \left| \frac{1}{\Delta t_2} \int_{t_1+\Delta t_s}^{t_2} \phi^1(t) dt - \frac{1}{\Delta t_1} \int_{t_0+\Delta t_s}^{t_1} \phi^1(t) dt \right| \right\|_2^2\end{aligned}$$

$$\begin{aligned}\Delta \mathbf{W}^1 &= -\eta_0 \frac{\partial L^1}{\partial \mathbf{W}^1} \\ \Delta \mathbf{W}^0 &= -\eta_1 \frac{\partial L^0}{\partial \mathbf{W}^0}\end{aligned}$$

Hidden layer target and loss

$$\phi_i^{0*} = \overline{\phi}_i^f + \alpha_i^t - \alpha_i^f$$

$$\begin{aligned}L^0 &\approx \|\phi^{0*} - \overline{\phi}^0\|_2^2 \\ &= \|\overline{\phi}^f + \alpha^t - \alpha_i^f - \overline{\phi}^0\|_2^2 \\ &= \|\alpha^t - \alpha^f\|_2^2\end{aligned}$$

*Actual loss functions more complicated and shown in methods

Defining the local loss functions

Output layer target and loss

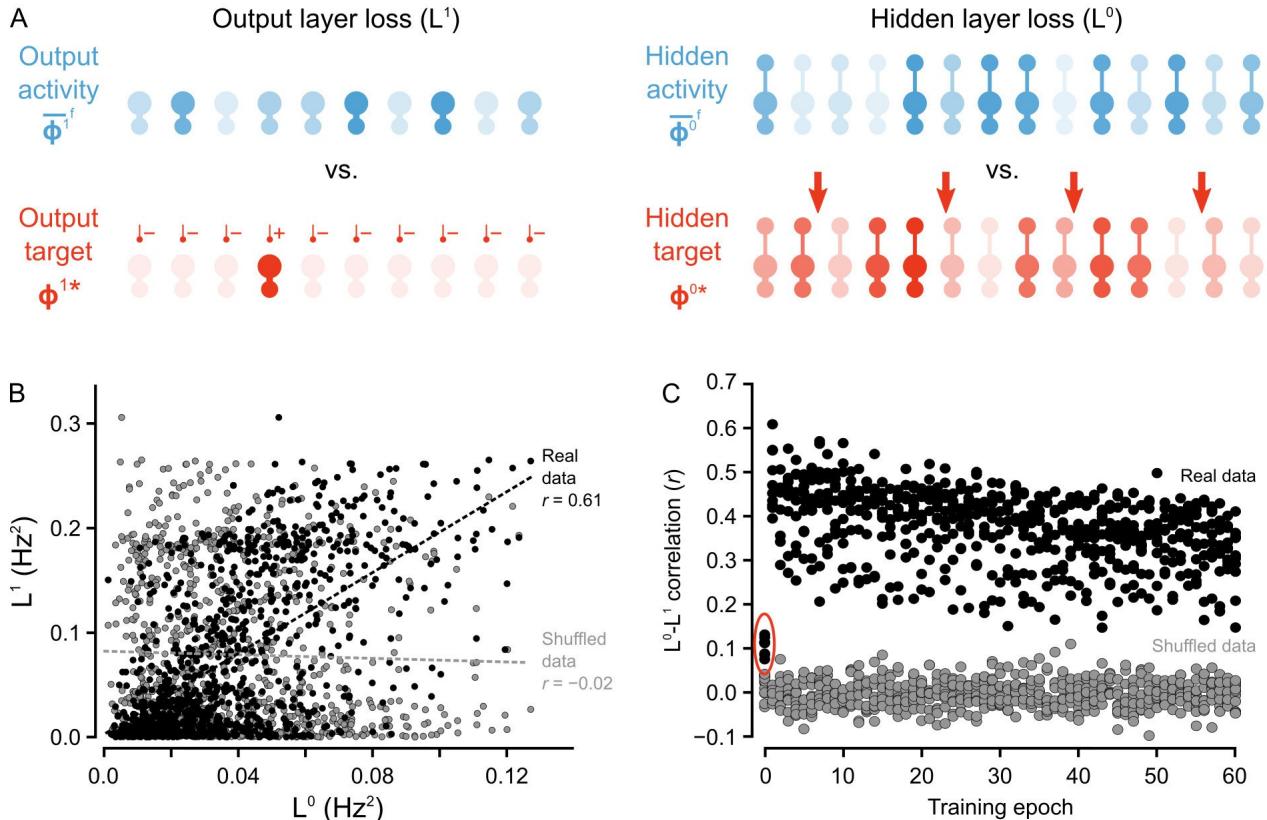
$$\begin{aligned}\phi_i^{1*} &= \overline{\phi_i^1}^t \\ &= \frac{1}{\Delta t_2} \int_{t_1+\Delta t_s}^{t_2} \phi_i^1(t) dt\end{aligned}$$

$$\begin{aligned}L^1 &\approx \|\phi^{1*} - \overline{\phi^1}^f\|_2^2 \\ &= \|\overline{\phi^1}^t - \overline{\phi^1}^f\|_2^2 \\ &= \left\| \left| \frac{1}{\Delta t_2} \int_{t_1+\Delta t_s}^{t_2} \phi^1(t) dt - \frac{1}{\Delta t_1} \int_{t_0+\Delta t_s}^{t_1} \phi^1(t) dt \right| \right\|_2^2\end{aligned}$$

Hidden layer target and loss

$$\phi_i^{0*} = \overline{\phi_i^0}^f + \alpha_i^t - \alpha_i^f$$

$$\begin{aligned}L^0 &\approx \|\phi^{0*} - \overline{\phi^0}^f\|_2^2 \\ &= \|\overline{\phi^0}^t + \alpha^t - \alpha_i^f - \overline{\phi^0}^f\|_2^2 \\ &= \|\alpha^t - \alpha^f\|_2^2\end{aligned}$$

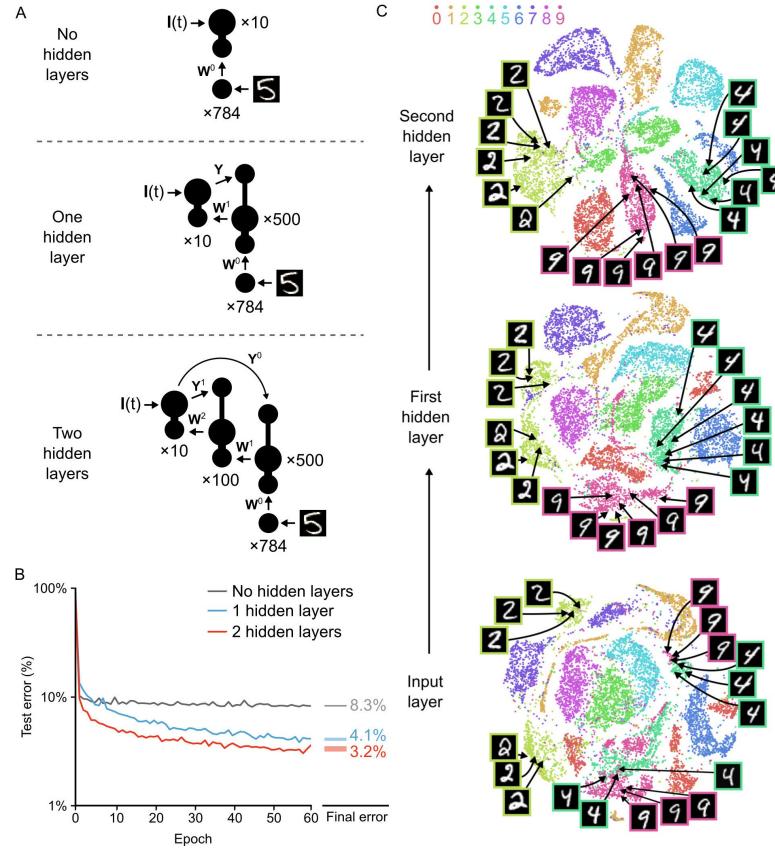


*Actual loss functions more complicated and shown in methods

Outline

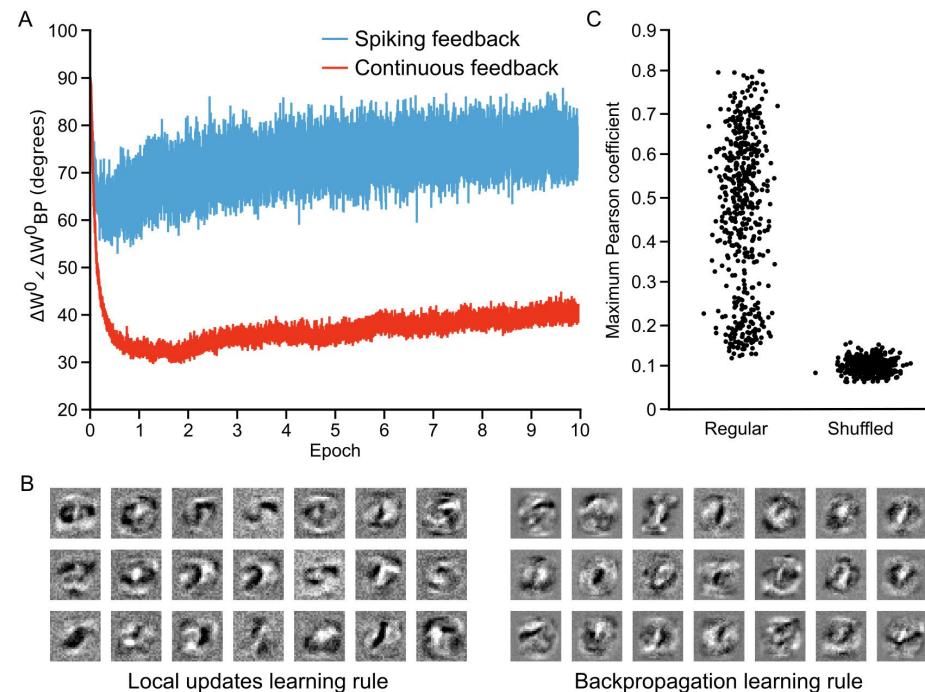
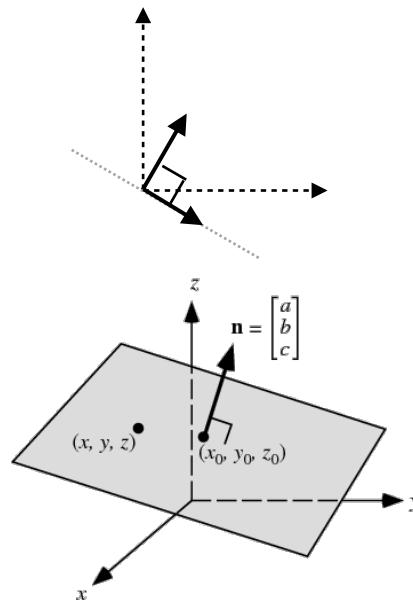
- Introduction: comparing the brain to a neural network
- The credit assignment problem in learning
- Primer on the machine learning solution: back-propagation
- Biology's solution? Deep learning with segregated dendrites
- **Results**
- Discussion

Deep learning: making good use of hidden layers



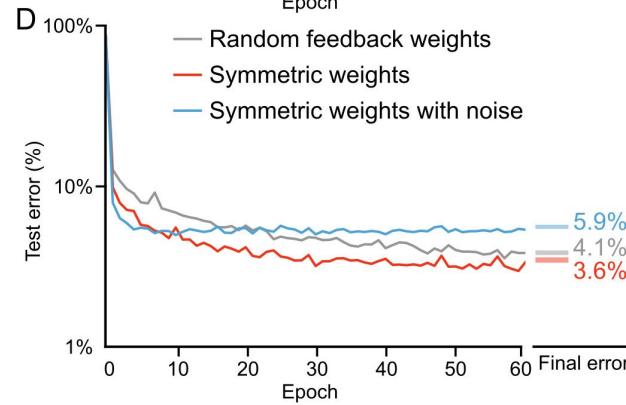
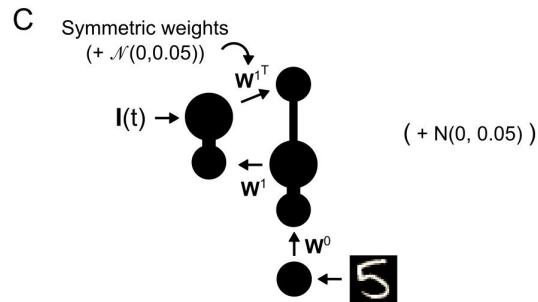
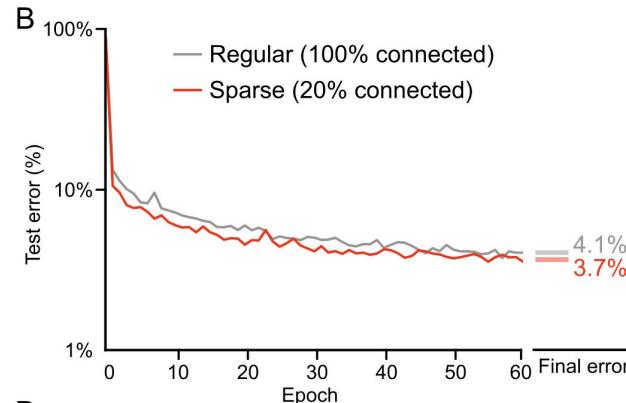
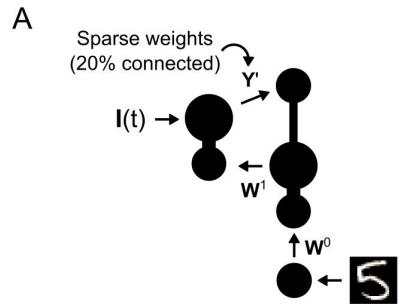
The algorithm approximates backpropagation in its weight updates

$$\Delta W^0 \text{ vs } \Delta W_{BP}^0$$

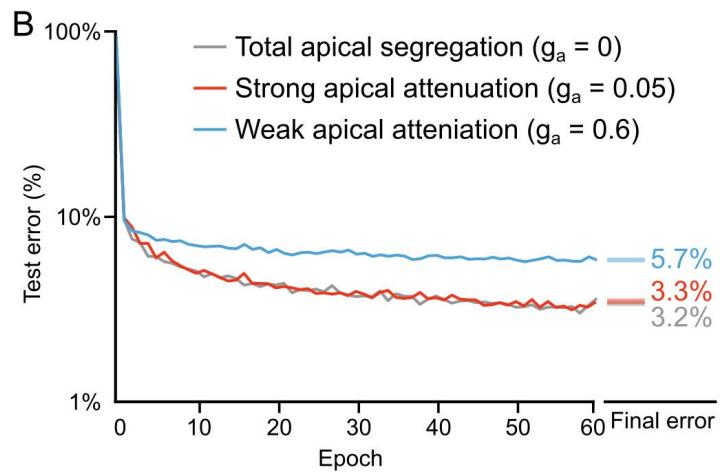
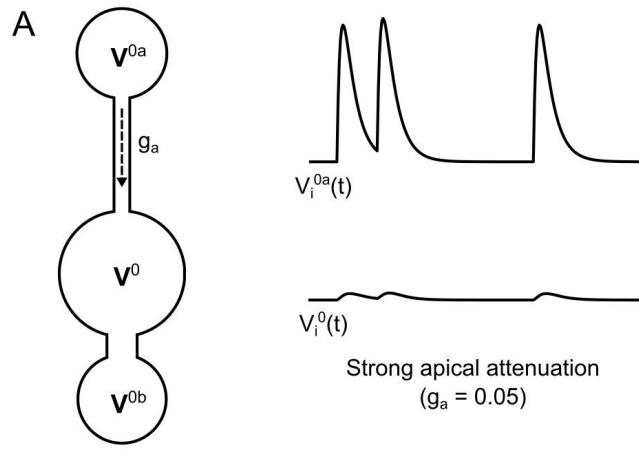


In higher dimensions, any random 2 vectors will be likely orthogonal

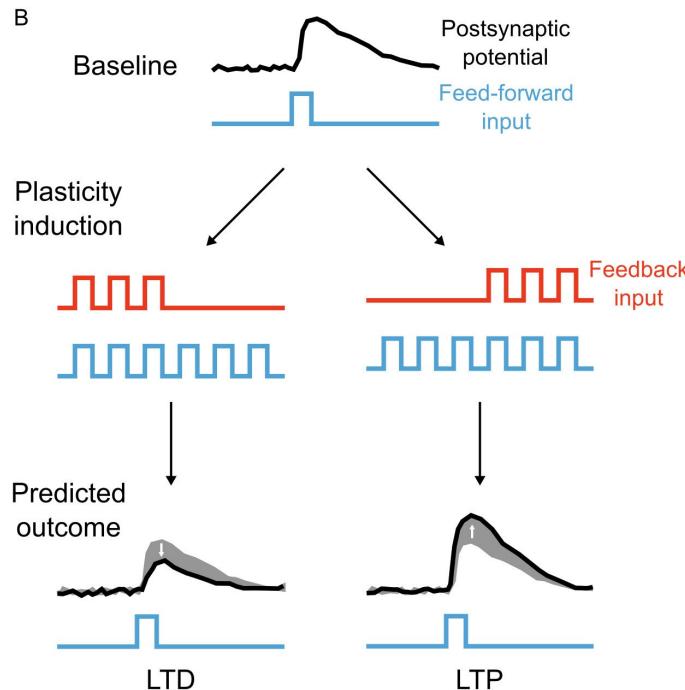
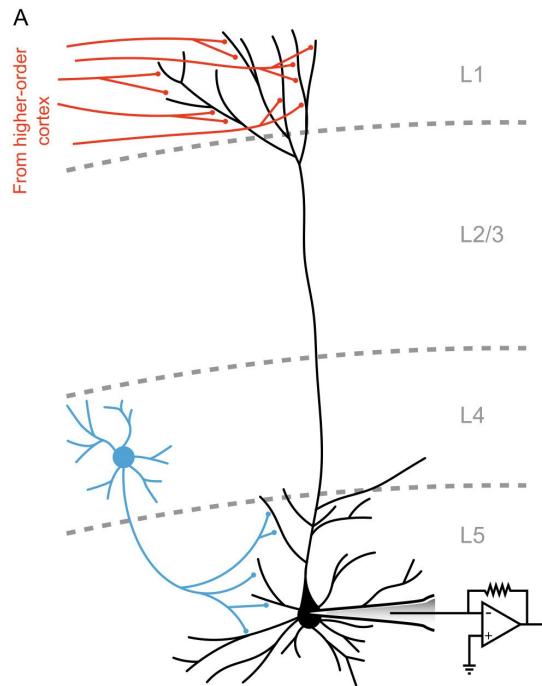
Conditions on feedback synapses



Apical dendritic segregation is important for learning



An experiment to test the central prediction of the model



Hidden layer target and loss

$$\phi_i^{0*} = \bar{\phi}_i^{0f} + \alpha_i^t - \alpha_i^f$$

$$\begin{aligned} L^0 &\approx \|\phi_i^{0*} - \bar{\phi}_i^{0f}\|_2^2 \\ &= \|\bar{\phi}_i^{0f} + \alpha_i^t - \alpha_i^f - \bar{\phi}_i^{0f}\|_2^2 \\ &= \|\alpha_i^t - \alpha_i^f\|_2^2 \end{aligned}$$

Plateau 1 higher than plateau 2:
LTD

Plateau 2 higher than plateau 1:
LTP

Summary

- Guergiev et al. solved the credit assignment problem with neurons with segregated dendrites
- Their algorithm approaches back-propagation in terms of the proposed weight updates
- Model of biologically realistic deep learning

Discussion

- Where does the teacher signal come from?
- Distinct phases of processing?
 - Might be governed by oscillations or neuromodulatory signals
- How do you overcome the temporal gap between plateau potentials?
 - Could be resolved by slower time constants

Thanks for listening!